

Well-Architected Framework (WAF) Review – DevOps CI/CD Project

Instructions

1. Baseline: Briefly describe your existing setup (tools, repos, target environment, runtime). 15 sentences max.
 2. Per pillar: Fill the sections below:
 - Current state
 - Gaps
 - TF improvements
 - Evidence
 3. Terraform: Implement or reference TF changes that realize improvements.
 4. Validation: Link evidence (TF code lines, logs, screenshots) where possible.
-

Baseline

- Tools: AWS (S3, CloudFront, API Gateway, Lambda, RDS, CloudWatch), CodePipeline, CodeBuild, SSM Parameter Store, Terraform, Go (tests), Node.js 18 (web, lambda), ESLint, Stylelint, Vitest.
 - Repositories/Monorepo: Single repo with **infra/** (Terraform), **cicd/** (pipelines), and **web/** (static site + lambda).
 - Environments (dev/stage/prod): Single "development" environment tags; pipelines branch-driven (**develop**).
 - Cloud provider/region(s): AWS **us-east-1** (set in buildspecs and Terraform backend).
 - Runtimes (frontend/backend/lambda/DB): Static web (HTML/CSS/JS), Node.js Lambda for contact form, PostgreSQL RDS, API Gateway REST API.
 - CI/CD (build, test, deploy): Two CodePipelines (infra & web) with CodeBuild steps; linting and tests in **buildspec-*.yaml**; S3 deploy, Lambda update, CloudFront invalidation.
 - IaC (Terraform versions/modules): Terraform **>=1.3** (backend S3/DynamoDB); modules for **s3**, **cloudfront**, **rds**, **lambda**, **api-gateway**, **iam**, **monitoring**, and **cicd**.
 - Observability (logs/metrics/tracing): CloudWatch logs for Lambda via **AWSLambdaBasicExecutionRole**; CloudWatch billing alarm; CodeBuild/CodePipeline logs.
 - Networking (VPCs/subnets): Uses default VPC only for RDS security group; Lambda not in VPC; API Gateway public; CloudFront over S3 OAI.
 - Security (IAM/KMS/secrets mgmt): IAM roles for CodeBuild/CodePipeline/Lambda; SSM parameters for DB creds; S3 public access blocked; CloudFront OAI for S3.
 - Data stores (RDS/S3/others): RDS Postgres instance; S3 website and artifacts buckets; SSM Parameter Store for config.
 - Edge/CDN: CloudFront distribution with HTTPS redirect; OAI restricting S3.
 - Key SLIs/SLOs: Not explicitly defined in repo.
-

1) Operational Excellence

Current state

- CI/CD via two CodePipelines with CodeBuild projects; infra pipeline runs fmt/validate/plan/apply; web pipeline runs lint/tests and deploys static site + Lambda + CF invalidation.
- Common tagging in Terraform locals; basic Go test scaffold present for infra.

Gaps

- No alarms/notifications on pipeline or build failures; no manual approval gates.
- Limited automated tests (Terratest skipped); no runbooks.

TF improvements

- Add CloudWatch alarms + SNS topics for CodeBuild/CodePipeline failures; add manual approval stage in pipelines.
- Add Terratest stage and enforce on PRs; expand tagging (owner, cost-center, service, environment).

Evidence

- `cicd/main.tf` (CodePipeline, CodeBuild stages)
 - `buildspec-infra.yml`, `buildspec-web.yml` (lint/test/plan/apply/deploy)
 - `infra/main.tf` locals `common_tags`
 - `infra/tests/infra_integration_test.go`
-

2) Security

Current state

- S3 website bucket blocks public access; CloudFront OAI policy restricts reads.
- API Gateway POST/OPTIONS without auth; RDS SG allows 0.0.0.0/0 to 5432 (demo).
- Terraform remote state in S3 with DynamoDB lock and encryption.
- DB credentials stored in SSM Parameter Store (`/rds/db_username`, `/rds/db_password` SecureString, `/rds/db_name`, `/rds/rds_address`) and read by Lambda at runtime.
- There are no Network ACLs (NACLs) configured for your VPC subnets. All subnet-level traffic filtering is handled by default AWS settings and security groups.

Gaps

- RDS publicly accessible SG; Lambda not in VPC; no KMS CMKs for S3/RDS.
- API lacks authentication/authorization and WAF; CodeBuild IAM policies broad with wildcards.
- No secret rotation; SSM path not scoped to environment.
- Missing stateless, subnet-level traffic filtering
- No defense-in-depth at the subnet layer
- Cannot explicitly deny unwanted IPs or ports at the subnet level

TF improvements

- Place RDS in private subnets and restrict SG to Lambda/VPC CIDR; attach KMS CMK to RDS and S3.
- Put Lambda in VPC with least-priv SG; add Secrets Manager for credentials with rotation; scope SSM paths by env.
- Add API auth (API key/JWT/Cognito) and AWS WAF ACL; tighten IAM to least privilege.
- Define a NACL resource in Terraform

Evidence

- `infra/modules/s3/main.tf` (public access block, OAI policy)
- `infra/modules/cloudfront/main.tf` (HTTPS redirect, OAI)
- `infra/modules/lambda/main.tf` (IAM role, SSM policy)
- `infra/modules/api-gateway/main.tf` (no auth)
- `infra/modules/rds/main.tf` (public SG 0.0.0.0/0)
- `infra/backend.tf` (S3 backend with DynamoDB lock)

Improvements made (code refs)

S3 website bucket versioning: `infra/modules/s3/main.tf` lines 9-15

```
resource "aws_s3_bucket_versioning" "website_versioning" {
  bucket = aws_s3_bucket.website.id
  versioning_configuration {
    status = "Enabled"
  }
}
```

S3 website bucket encryption: `infra/modules/s3/main.tf` lines 17-26

```
resource "aws_s3_bucket_server_side_encryption_configuration"
"website_encryption" {
  bucket = aws_s3_bucket.website.id
  rule { apply_server_side_encryption_by_default { sse_algorithm =
"AES256" } }
}
```

S3 artifacts versioning: `infra/modules/s3/main.tf` lines 79-84

```
resource "aws_s3_bucket_versioning" "codepipeline_artifacts_versioning"
{
  bucket = aws_s3_bucket.codepipeline_artifacts.id
  versioning_configuration {
    status = "Enabled"
  }
}
```

S3 artifacts encryption: [infra/modules/s3/main.tf](#) lines 86-95

```
resource "aws_s3_bucket_server_side_encryption_configuration"
"codepipeline_artifacts_encryption" {
  bucket = aws_s3_bucket.codepipeline_artifacts.id

  rule {
    apply_server_side_encryption_by_default {
      sse_algorithm = "AES256"
    }
  }
}
```

CloudFront security headers policy: [infra/modules/cloudfront/main.tf](#) lines 20-37

```
default_cache_behavior {
  allowed_methods  = ["GET", "HEAD"]
  cached_methods  = ["GET", "HEAD"]
  target_origin_id = "s3-origin"

  forwarded_values {
    query_string = false
    cookies {
      forward = "none"
    }
  }
}

viewer_protocol_policy    = "redirect-to-https"
min_ttl                   = 0
default_ttl                = 3600
max_ttl                   = 86400
response_headers_policy_id = "60669652-455b-4ae9-85a4-c4c02393f86c" #
AWSManagedSecurityHeadersPolicy
}
```

CloudFront minimum TLS version: [infra/modules/cloudfront/main.tf](#) lines 45-48

```
viewer_certificate {
  cloudfront_default_certificate = true
  minimum_protocol_version       = "TLSv1.2_2021"
}
```

API Gateway access logs: [infra/modules/api-gateway/main.tf](#) lines 101-123

```

resource "aws_api_gateway_stage" "contact_stage" {
  deployment_id = aws_api_gateway_deployment.contact_deployment.id
  rest_api_id   = aws_api_gateway_rest_api.contact_api.id
  stage_name    = var.stage_name

  access_log_settings {
    destination_arn = aws_cloudwatch_log_group.api_gw_logs.arn
    format = jsonencode({
      requestId      = "$context.requestId",
      ip             = "$context.identity.sourceIp",
      caller         = "$context.identity.caller",
      user           = "$context.identity.user",
      requestTime    = "$context.requestTime",
      httpMethod     = "$context.httpMethod",
      resourcePath   = "$context.resourcePath",
      status         = "$context.status",
      protocol       = "$context.protocol",
      responseLength = "$context.responseLength",
      integrationStatus = "$context.integration.status",
      integrationError = "$context.integrationErrorMessage"
    })
  }
}

```

API Gateway log group and method throttling: [infra/modules/api-gateway/main.tf](#) lines 125-145

```

resource "aws_cloudwatch_log_group" "api_gw_logs" {
  name =
"/apigw/${aws_api_gateway_rest_api.contact_api.id}/${var.stage_name}"
  retention_in_days = 14
  tags = var.tags
}

resource "aws_api_gateway_method_settings" "all" {
  rest_api_id = aws_api_gateway_rest_api.contact_api.id
  stage_name  = aws_api_gateway_stage.contact_stage.stage_name
  method_path = "/*/*"

  settings {
    metrics_enabled      = true
    logging_level        = "INFO"
    data_trace_enabled   = true
    throttling_burst_limit = 5
    throttling_rate_limit = 10
  }
}

```

```
resource "aws_wafv2_web_acl" "apigw_acl" {
  name          = "apigw-basic-acl"
  description   = "Basic protections for API Gateway"
  scope         = "REGIONAL"
  default_action {
    allow {}
  }
  rule {
    name          = "AWS-AWSManagedRulesCommonRuleSet"
    priority      = 1
    statement {
      managed_rule_group_statement {
        name          = "AWSManagedRulesCommonRuleSet"
        vendor_name   = "AWS"
      }
    }
    visibility_config {
      cloudwatch_metrics_enabled = true
      metric_name                 = "AWSManagedRulesCommonRuleSet"
      sampled_requests_enabled    = true
    }
  }
  visibility_config {
    cloudwatch_metrics_enabled = true
    metric_name                 = "apigw-acl"
    sampled_requests_enabled    = true
  }
  tags = var.tags
}

resource "aws_wafv2_web_acl_association" "apigw_acl_assoc" {
  resource_arn = aws_api_gateway_stage.contact_stage.arn
  web_acl_arn  = aws_wafv2_web_acl.apigw_acl.arn
}
```

```
resource "aws_iam_role_policy_attachment" "lambda_vpc_access" {
  role          = aws_iam_role.lambda_exec.name
  policy_arn    = "arn:aws:iam::aws:policy/service-role/AWSLambdaVPCEAccessExecutionRole"
}
```

```

resource "aws_security_group" "lambda_sg" {
  name_prefix = "lambda-sg-"
  description = "Security group for Lambda to access RDS"
  vpc_id      = data.aws_vpc.default.id

  egress {
    from_port = 0
    to_port   = 0
    protocol  = "-1"
    cidr_blocks = ["0.0.0.0/0"]
  }

  tags = merge(var.tags, { Name = "lambda-sg" })
}

```

Lambda VPC config on function: [infra/modules/lambda/main.tf](#) lines 89-92

```

vpc_config {
  subnet_ids          = data.aws_subnets.default_vpc_subnets.ids
  security_group_ids = [aws_security_group.lambda_sg.id]
}

```

RDS SG ingress from Lambda SG: [infra/modules/rds/main.tf](#) lines 6-17

```

resource "aws_security_group" "rds_ingress" {
  name_prefix = "rds-ingress-5432-"
  description = "Allow inbound to Postgres from Lambda SG"
  vpc_id      = data.aws_vpc.default.id

  ingress {
    from_port = 5432
    to_port   = 5432
    protocol  = "tcp"
    security_groups = [var.allowed_sg_id]
    description    = "Allow Postgres from Lambda security
group"
  }
}

```

RDS private/encrypted + SG attach: [infra/modules/rds/main.tf](#) lines 48-58

```

resource "aws_db_instance" "contact_db" {
  storage_encrypted = var.storage_encrypted
  publicly_accessible = var.publicly_accessible
}

```

```
vpc_security_group_ids = [aws_security_group.rds_ingress.id]
}
```

IAM narrowed S3 bucket resources: [infra/modules/iam/main.tf](#) lines 276-279

```
Resource = [
    var.artifacts_bucket_arn,
    var.website_bucket_arn
]
```

Secrets Manager secret and version: [infra/main.tf](#) lines 58-75

```
resource "aws_secretsmanager_secret" "db_credentials" {
    name          = "project3/db-credentials"
    description    = "Database credentials for contact form"
    tags          = local.common_tags
}

resource "aws_secretsmanager_secret_version" "db_credentials_version" {
    secret_id      = aws_secretsmanager_secret.db_credentials.id
    secret_string  = jsonencode({
        username = coalesce(var.db_username,
data.aws_ssm_parameter.db_username.value)
        password = coalesce(var.db_password,
data.aws_ssm_parameter.db_password.value)
        host      = module.rds.rds_address
        database  = coalesce(var.db_name,
data.aws_ssm_parameter.db_name.value)
        port      = module.rds.rds_port
    })
    depends_on    = [module.rds]
}
```

- Rotation function and rotation rule: [infra/main.tf](#) lines 76-103

```
resource "aws_serverlessapplicationrepository_cloudformation_stack"
"rds_rotation" {
    name          = "project3-rds-rotation"
    application_id = "arn:aws:serverlessrepo:us-east-
1:297356227824:applications/SecretsManagerRDSPostgreSQLRotationSingleUse
r"
    capabilities  = ["CAPABILITY_NAMED_IAM"]
    parameters = {
        functionName      = "project3-rds-rotation"
        vpcSubnetIds       = join(",",
data.aws_subnets.default_vpc_subnets.ids)
    }
}
```



```

    vpcSecurityGroupIds = module.lambda.lambda_security_group_id
  }
  semantic_version = "1.1.188"
  tags             = local.common_tags
}

resource "aws_secretsmanager_secret_rotation" "db_rotation" {
  secret_id          = aws_secretsmanager_secret.db_credentials.id
  rotation_lambda_arn =
aws_serverlessapplicationrepository_cloudformation_stack.rds_rotation.ou
tputs["RotationLambdaARN"]
  rotation_rules { automatically_after_days = 30 }
  depends_on =
[aws_secretsmanager_secret_version.db_credentials_version]
}

```

- Lambda environment + IAM for secret read: [infra/modules/lambda/main.tf](#) lines 56-66, 70-90

```

environment {
  variables = {
    ENVIRONMENT = "development"
    DB_SECRET_ARN = var.db_secret_arn
  }
}

```

```

resource "aws_iam_role_policy" "lambda_secrets_policy" {
  name = "lambda-secrets-access"
  role = aws_iam_role.lambda_exec.id
  policy = jsonencode({
    Version = "2012-10-17",
    Statement = [{
      Effect = "Allow",
      Action =
["secretsmanager:GetSecretValue","secretsmanager:DescribeSecret"],
      Resource = var.db_secret_arn
    }]
  })
}

```

- Lambda code reads from Secrets Manager with SSM fallback: [web/lambda/index.js](#) lines 1-18, 21-40, 44-74

```

// Security note: Database credentials are retrieved from AWS SSM
Parameter Store, not environment variables.
import { Client } from "pg";

```

```
import { SSMClient, GetParameterCommand } from "@aws-sdk/client-ssm";
import { SecretsManagerClient, GetSecretValueCommand } from "@aws-
sdk/client-secrets-manager";

// Initialize AWS clients
const region = process.env.AWS_REGION || 'us-east-1';
const ssmClient = new SSMClient({ region });
const secretsClient = new SecretsManagerClient({ region });
```

```
// Cache for database credentials to avoid repeated SSM calls
let dbCredentials = null;

// Function to get database credentials (Secrets Manager preferred,
// fallback to SSM)
async function getDbCredentials() {
  if (dbCredentials) {
    return dbCredentials;
  }

  try {
    const secretArn = process.env.DB_SECRET_ARN;
    if (secretArn) {
      const secretData = await secretsClient.send(new
GetSecretValueCommand({ SecretId: secretArn }));
      const secret = JSON.parse(secretData.SecretString || '{}');
      dbCredentials = { host: secret.host, user: secret.username,
password: secret.password, database: secret.database };
    } else {
      // Fallback to SSM parameters
```

```
const [dbHost, dbUser, dbPass, dbName] = await Promise.all([
  ssmClient.send(new GetParameterCommand({ Name:
"/rds/rds_address" })),
  ssmClient.send(new GetParameterCommand({ Name:
"/rds/db_username" })),
  ssmClient.send(new GetParameterCommand({ Name:
"/rds/db_password", WithDecryption: true })),
  ssmClient.send(new GetParameterCommand({ Name: "/rds/db_name"
}))
]);
dbCredentials = { host: dbHost.Parameter.Value, user:
dbUser.Parameter.Value, password: dbPass.Parameter.Value, database:
dbName.Parameter.Value };
return dbCredentials;
} catch (error) {
  console.error("Failed to retrieve database credentials:", error);
  throw new Error("Database configuration error");
}
```

```
}  
}
```

-Network ACLs (NACLs) have been implemented in your Terraform VPC module.[infra/vpc/main.tf](#)
line 124-194

```
# Network ACLs for Public Subnets  
resource "aws_network_acl" "public" {  
  vpc_id = aws_vpc.main.id  
  tags = merge(var.tags, { Name = "${var.environment}-public-nacl" })  
}  
  
# Allow HTTPS inbound, deny all else (example)  
resource "aws_network_acl_rule" "public_https_inbound" {  
  network_acl_id = aws_network_acl.public.id  
  rule_number    = 100  
  egress         = false  
  protocol       = "tcp"  
  rule_action    = "allow"  
  cidr_block     = "0.0.0.0/0"  
  from_port      = 443  
  to_port        = 443  
}  
  
resource "aws_network_acl_rule" "public_deny_all_inbound" {  
  network_acl_id = aws_network_acl.public.id  
  rule_number    = 200  
  egress         = false  
  protocol       = "-1"  
  rule_action    = "deny"  
  cidr_block     = "0.0.0.0/0"  
  from_port      = 0  
  to_port        = 0  
}  
  
# Associate Public NACL with Public Subnets  
resource "aws_network_acl_association" "public" {  
  count          = length(var.public_subnet_cidrs)  
  subnet_id     = aws_subnet.public[count.index].id  
  network_acl_id = aws_network_acl.public.id  
}  
  
# Network ACLs for Private Subnets  
resource "aws_network_acl" "private" {  
  vpc_id = aws_vpc.main.id  
  tags = merge(var.tags, { Name = "${var.environment}-private-nacl" })  
}  
  
# Allow DB traffic from Lambda SG CIDR (example: adjust as needed)  
resource "aws_network_acl_rule" "private_db_inbound" {
```

```

    network_acl_id = aws_network_acl.private.id
    rule_number    = 100
    egress         = false
    protocol       = "tcp"
    rule_action    = "allow"
    cidr_block     = "10.0.0.0/8" # Example CIDR, adjust to Lambda SG
  subnet
    from_port     = 5432
    to_port       = 5432
}

resource "aws_network_acl_rule" "private_deny_all_inbound" {
  network_acl_id = aws_network_acl.private.id
  rule_number    = 200
  egress         = false
  protocol       = "-1"
  rule_action    = "deny"
  cidr_block     = "0.0.0.0/0"
  from_port      = 0
  to_port        = 0
}

# Associate Private NACL with Private Subnets
resource "aws_network_acl_association" "private" {
  count          = length(var.private_subnet_cidrs)
  subnet_id     = aws_subnet.private[count.index].id
  network_acl_id = aws_network_acl.private.id
}

```

3) Reliability

Current state

- Single RDS instance in primary region without Multi-AZ
- Basic RDS backups and maintenance windows configured
- No cross-region disaster recovery setup
- No automated failover mechanism
- RPO/RTO requirements not met (e-commerce needs: RPO 1h, RTO 4h)

Gaps

- Single point of failure with non-Multi-AZ RDS
- No cross-region redundancy for disaster recovery
- Missing health checks and automated failover
- No warm standby setup to meet RTO requirement
- Backup strategy insufficient for RPO requirement
- No DLQ or retries on Lambda functions
- Missing monitoring and alerting system

TF improvements

1. Implement Warm Standby Architecture:

- Deploy standby RDS in us-west-2
- Configure cross-region replication
- Set up Route53 health checks and failover routing
- Enable Multi-AZ for primary RDS

2. Enhance Monitoring and Recovery:

- Add CloudWatch alarms for API errors (4XX/5XX)
- Monitor Lambda performance and failures
- Create operational dashboards
- Implement automated failover testing

3. Improve Data Protection:

- Configure RDS automated backups every 15 minutes
- Enable point-in-time recovery
- Implement cross-region S3 replication
- Set up proper backup retention policies

Evidence

- `infra/modules/monitoring/main.tf` (only billing alarm)
- `infra/modules/api-gateway/main.tf` (stage definition)
- `infra/modules/lambda/main.tf` (no DLQ)
- `infra/modules/rds/main.tf` (backup/deletion flags)

Reliability improvements (code refs)

Warm Standby Architecture Decision

The decision to implement a warm standby architecture was based on the following requirements and considerations:

1. Recovery Time Objective (RTO):

- Requirement: 4 hours maximum downtime allowed
- Warm standby provides faster recovery compared to cold standby or backup/restore
- Pre-provisioned infrastructure reduces deployment time during failover

2. Recovery Point Objective (RPO):

- Requirement: Maximum 1 hour of data loss acceptable
- Continuous replication of data to standby region meets this requirement
- S3 cross-region replication for static content
- Database replication for dynamic data

3. Traffic Pattern Analysis:

- Steady, predictable traffic pattern
- Non-spiky workload suits warm standby's cost-effectiveness
- Lower cost compared to active-active while meeting RPO/RTO

4. Cost-Benefit Analysis:

- Warm standby provides optimal balance between recovery speed and cost
- Standby resources can run on smaller instances to reduce costs
- No need for complex active-active synchronization

5. Operational Complexity:

- Simpler than active-active architecture
- Automated failover through Route53 health checks
- Clear, well-defined failover process
- Easier to test and maintain

Summary of Reliability Improvements

We implemented a warm standby architecture and enhanced reliability with the following changes:

- **Created standby RDS module in a different region:**
 - Folder: `infra/modules/rds-standby/`
 - Files: `main.tf`, `variables.tf`, `outputs.tf`
 - Purpose: Deploys a scaled-down standby PostgreSQL RDS instance in `us-west-2`.
- **Enabled S3 cross-region replication for static assets:**
 - File: `infra/modules/s3/replication.tf`
 - Purpose: Replicates website bucket to standby region for DR.
- **Configured Route53 DNS failover for APIs:**
 - File: `infra/modules/route53/failover.tf`
 - Supporting files: `variables.tf`, `outputs.tf`
 - Purpose: Automated DNS failover between primary and standby API endpoints.
- **Documented the warm standby setup and failover process:**
 - File: `WARM-STANDBY-README.md` (project root)
 - Purpose: Instructions and validation steps for disaster recovery.

All new modules follow Terraform best practices with variables, outputs, and clear separation of primary/standby resources.

Key code implementations:

1. RDS Standby Configuration:

```
# infra/modules/rds-standby/main.tf
resource "aws_db_instance" "contact_db_standby" {
  identifier          = "contact-db-standby"
  engine              = "postgres"
  instance_class      = "db.t3.micro" # Scaled down for cost in standby
  allocated_storage   = 20
  storage_encrypted   = true

  backup_retention_period = 7
  backup_window           = "03:00-04:00"
  maintenance_window     = "Mon:04:00-Mon:05:00"

  multi_az            = false # Single AZ for standby to reduce costs
  publicly_accessible = false

  vpc_security_group_ids = [aws_security_group.rds_standby_sg.id]
  db_subnet_group_name   = aws_db_subnet_group.standby_subnet_group.name

  tags = merge(var.tags, {
    Name = "contact-db-standby"
    Role = "warm-standby"
  })
}
```

2. S3 Cross-Region Replication:

```
# infra/modules/s3/replication.tf
resource "aws_s3_bucket_replication_configuration" "website_replication"
{
  role = aws_iam_role.replication_role.arn
  bucket = aws_s3_bucket.website.id

  rule {
    id      = "website-standby-replication"
    status = "Enabled"

    destination {
      bucket      = aws_s3_bucket.website_standby.arn
      storage_class = "STANDARD"
    }
  }
}
```

3. Route53 DNS Failover:

```
# infra/modules/route53/failover.tf
resource "aws_route53_health_check" "primary_api" {
  fqdn = var.primary_api_domain
}
```

```

port          = 443
type          = "HTTPS"
resource_path  = "/health"
failure_threshold = "3"
request_interval = "30"

tags = merge(var.tags, {
  Name = "primary-api-health-check"
})
}

resource "aws_route53_record" "api_primary" {
  zone_id = var.hosted_zone_id
  name     = var.api_domain
  type     = "A"

  failover_routing_policy {
    type = "PRIMARY"
  }

  alias {
    name           = var.primary_api_domain
    zone_id        = var.primary_api_zone_id
    evaluate_target_health = true
  }

  health_check_id = aws_route53_health_check.primary_api.id
  set_identifier   = "primary"
}

resource "aws_route53_record" "api_secondary" {
  zone_id = var.hosted_zone_id
  name     = var.api_domain
  type     = "A"

  failover_routing_policy {
    type = "SECONDARY"
  }

  alias {
    name           = var.standby_api_domain
    zone_id        = var.standby_api_zone_id
    evaluate_target_health = true
  }

  set_identifier = "secondary"
}

---

- Lambda reserved concurrency: `infra/modules/lambda/main.tf` lines 109-
111
```109:111:infra/modules/lambda/main.tf

```



```
Reserve concurrency to prevent thundering herds and protect DB
reserved_concurrent_executions = 5
```

- Lambda DLQ and invoke config: [infra/modules/lambda/main.tf](#) lines 130-145

```
resource "aws_sqs_queue" "lambda_dlq" {
 name = "contact-form-dlq"
 tags = var.tags
}

resource "aws_lambda_event_invoke_config" "contact_eic" {
 function_name = aws_lambda_function.contact.function_name
 destination_config { on_failure { destination =
aws_sqs_queue.lambda_dlq.arn } }
 maximum_retry_attempts = 2
 maximum_event_age_in_seconds = 3600
}
```

- Lambda errors alarm: [infra/modules/monitoring/main.tf](#) lines 25-42

```
resource "aws_cloudwatch_metric_alarm" "lambda_errors" {
 alarm_name = "lambda-contact-errors"
 comparison_operator = "GreaterThanThreshold"
 evaluation_periods = 1
 metric_name = "Errors"
 namespace = "AWS/Lambda"
 period = 60
 statistic = "Sum"
 threshold = 1
 alarm_actions = var.alarm_actions
 dimensions = { FunctionName = var.lambda_function_name }
}
```

- Lambda p95 duration alarm: [infra/modules/monitoring/main.tf](#) lines 44-60

```
resource "aws_cloudwatch_metric_alarm" "lambda_duration" {
 alarm_name = "lambda-contact-duration-p95"
 comparison_operator = "GreaterThanThreshold"
 evaluation_periods = 1
 metric_name = "Duration"
 namespace = "AWS/Lambda"
 period = 60
 statistic = "p95"
 threshold = 3000
 alarm_actions = var.alarm_actions
}
```

```

 dimensions = { FunctionName = var.lambda_function_name }
}

```

- API Gateway 5XX alarm: [infra/modules/monitoring/main.tf](#) lines 62-79

```

resource "aws_cloudwatch_metric_alarm" "apigw_5xx" {
 alarm_name = "apigw-5xx-errors"
 comparison_operator = "GreaterThanThreshold"
 evaluation_periods = 1
 metric_name = "5XXError"
 namespace = "AWS/ApiGateway"
 period = 60
 statistic = "Sum"
 threshold = 1
 alarm_actions = var.alarm_actions
 dimensions = { ApiId = var.api_gateway_id, Stage =
var.api_gateway_stage, Resource = "/contact", Method = "POST" }
}

```

- Multi-AZ for primary RDS and DMS cross-region replication:

```

resource "aws_db_instance" "contact_db" {
 # ...existing config...
 multi_az = true
 # Cross-region replication with AWS DMS
 # See DMS resources below
}

```

- DMS replication resources:

```

resource "aws_dms_replication_instance" "rds_replication" { ... }
resource "aws_dms_endpoint" "source" { ... }
resource "aws_dms_endpoint" "target" { ... }
resource "aws_dms_replication_task" "rds_to_standby" { ... }
resource "aws_dms_replication_subnet_group" "dms_subnet_group" { ... }

```

- DMS table mappings and task settings:

```

{
 "rules": [
 {
 "rule-type": "selection",
 "rule-id": "1",
 "rule-name": "1",

```

```

 "object-locator": {
 "schema-name": "%",
 "table-name": "%"
 },
 "rule-action": "include"
 }
]
}

```

```

{
 "TargetMetadata": { ... },
 "FullLoadSettings": { ... },
 "Logging": { ... },
 "ControlTablesSettings": { ... },
 "StreamBufferSettings": { ... },
 "ChangeProcessingDdlHandlingPolicy": { ... },
 "ErrorBehavior": { ... },
 "ChangeProcessingPolicy": { ... }
}

```

- S3 lifecycle for backup retention:

```

resource "aws_s3_bucket_lifecycle_configuration" "website_lifecycle" {
 bucket = aws_s3_bucket.website.id
 rule {
 id = "cleanup_old_versions"
 status = "Enabled"
 filter {
 prefix = ""
 }
 noncurrent_version_expiration {
 noncurrent_days = 30
 }
 }
}

```

---

## 4) Performance Efficiency

### Current state

Initial performance analysis revealed several areas needing optimization:

1. Content Delivery Network (CDN):
  - S3 + CloudFront distribution with basic 24-hour TTL
  - No compression enabled for static assets

- Missing Cache-Control headers for browser caching
- Default CloudFront settings without optimization

## 2. API Gateway and Lambda:

- Default Lambda configuration (128MB memory, 3s timeout)
- No API Gateway caching implemented
- Missing throttling and concurrency controls
- Basic database connection handling

## 3. Database Layer:

- RDS PostgreSQL on db.t3.micro (baseline tier)
- Default parameter group without tuning
- No Performance Insights monitoring
- Basic connection management

## 4. Performance Monitoring:

- Limited visibility into system performance
- No comprehensive monitoring dashboard
- Missing performance metrics collection
- No automated performance alerting

# Improvements Made

## 1. Content Delivery Network Optimization:

- Enabled CloudFront compression with Brotli support
- Implemented tiered TTL strategy:
  - Static assets: 1-year cache (31536000s)
  - Dynamic content: 24-hour cache (86400s)
  - API responses: 5-minute cache (300s)
- Added Cache-Control headers via CloudFront policies
- Configured optimal compression settings for asset types

## 2. API and Lambda Performance Enhancements:

- Increased Lambda memory to 256MB for better performance
- Extended timeout to 10s for reliable operation
- Added API Gateway caching (0.5GB cache size)
- Implemented database connection pooling
- Set up provisioned concurrency (2 instances)

## 3. Database Optimizations:

- Upgraded to db.t3.small for improved performance
- Enabled Performance Insights (7-day retention)
- Optimized database parameters:
  - Shared buffers: 20% of instance memory

- Work memory: 8MB per connection
- Max connections: 100
- Implemented enhanced monitoring (60s intervals)

#### 4. Performance Monitoring Framework:

- Created comprehensive CloudWatch dashboard
- Set up performance metric alarms
- Added API Gateway metrics tracking
- Implemented Lambda performance monitoring

### Evidence

- CloudFront basic config: [infra/modules/cloudfront/main.tf](#) lines 15-25:

```
default_cache_behavior {
 min_ttl = 0
 default_ttl = 86400 # 24 hours default
 max_ttl = 31536000 # 1 year max
}
```

- Lambda basic settings: [infra/modules/lambda/main.tf](#) lines 45-48:

```
resource "aws_lambda_function" "contact_form" {
 memory_size = 128 # Default memory
 timeout = 3 # Default timeout
}
```

- API Gateway without caching: [infra/modules/api-gateway/main.tf](#) lines 89-95:

```
resource "aws_api_gateway_stage" "api" {
 # No cache_cluster_enabled setting
 # No method_settings for caching
}
```

- RDS base configuration: [infra/modules/rds/main.tf](#) lines 25-30:

```
resource "aws_db_instance" "contact_db" {
 instance_class = "db.t3.micro"
 # No performance_insights_enabled
 # Default parameter group
}
```

## 1. CloudFront Optimizations

CloudFront compression and caching: [infra/modules/cloudfront/main.tf](#) lines 30-52

```
resource "aws_cloudfront_distribution" "website" {
 # ... existing configuration ...

 default_cache_behavior {
 compress = true
 viewer_protocol_policy = "redirect-to-https"

 min_ttl = 0
 default_ttl = 86400 # 24 hours
 max_ttl = 31536000 # 1 year

 cached_methods = ["GET", "HEAD", "OPTIONS"]
 allowed_methods = ["GET", "HEAD", "OPTIONS"]

 forwarded_values {
 query_string = false
 cookies { forward = "none" }
 headers = ["Origin", "Access-Control-Request-Method",
"Access-Control-Request-Headers"]
 }
 }

 ordered_cache_behavior {
 path_pattern = "/assets/*"
 allowed_methods = ["GET", "HEAD"]
 cached_methods = ["GET", "HEAD"]
 target_origin_id = aws_s3_bucket.website.id

 compress = true
 viewer_protocol_policy = "redirect-to-https"

 min_ttl = 86400 # 1 day
 default_ttl = 604800 # 1 week
 max_ttl = 31536000 # 1 year
 }
}
```

CloudFront response headers policy: [infra/modules/cloudfront/main.tf](#) lines 54-75

```
resource "aws_cloudfront_response_headers_policy" "security_headers" {
 name = "performance-security-headers"

 custom_headers_config {
```

```

 items {
 header = "Cache-Control"
 value = "public, max-age=31536000"
 override = true
 }
 items {
 header = "Accept-Encoding"
 value = "gzip, deflate, br"
 override = true
 }
 }

 security_headers_config {
 content_type_options {
 override = true
 }
 frame_options {
 frame_option = "DENY"
 override = true
 }
 strict_transport_security {
 access_control_max_age_sec = 31536000
 include_subdomains = true
 override = true
 }
 }
}

```

## 2. API Gateway Caching

API Gateway cache settings: [infra/modules/api-gateway/main.tf](#) lines 150-180

```

resource "aws_api_gateway_stage" "contact_stage" {
 # ... existing configuration ...

 cache_cluster_enabled = true
 cache_cluster_size = "0.5" # 0.5GB cache

 variables = {
 "cacheEnabled" = "true"
 }
}

resource "aws_api_gateway_method_settings" "contact_cache" {
 rest_api_id = aws_api_gateway_rest_api.contact_api.id
 stage_name = aws_api_gateway_stage.contact_stage.stage_name
 method_path = "*/*"

 settings {
 metrics_enabled = true
 }
}

```

```

 logging_level = "INFO"
 caching_enabled = true
 cache_ttl_in_secs = 300 # 5 minutes cache

 throttling_burst_limit = 100
 throttling_rate_limit = 50
 }
}

```

### 3. Lambda Optimization

Lambda performance configuration: [infra/modules/lambda/main.tf](#) lines 80-110

```

resource "aws_lambda_function" "contact_form" {
 # ... existing configuration ...

 memory_size = 256
 timeout = 10

 environment {
 variables = {
 NODE_OPTIONS = "--enable-source-maps"
 POSTGRES_MAX_CONNECTIONS = "10"
 }
 }

 provisioned_concurrent_executions = 2

 vpc_config {
 subnet_ids = data.aws_subnets.default_vpc_subnets.ids
 security_group_ids = [aws_security_group.lambda_sg.id]
 }
}

Add connection pooling utility
resource "aws_lambda_layer_version" "pg_pool" {
 layer_name = "pg-pool-layer"
 description = "PostgreSQL connection pooling layer"
 filename = "layers/pg-pool.zip"
 compatible_runtimes = ["nodejs18.x"]
}

```

### 4. RDS Performance

RDS performance configuration: [infra/modules/rds/main.tf](#) lines 85-120

```

resource "aws_db_parameter_group" "contact_db_params" {
 name = "contact-db-params"
}

```



```

family = "postgres14"

parameter {
 name = "shared_buffers"
 value = "{DBInstanceClassMemory*20/100}" # 20% of instance memory
}

parameter {
 name = "max_connections"
 value = "100"
}

parameter {
 name = "work_mem"
 value = "8388608" # 8MB
}
}

resource "aws_db_instance" "contact_db" {
 # ... existing configuration ...

 instance_class = "db.t3.small"

 # Enable performance insights
 performance_insights_enabled = true
 performance_insights_retention_period = 7
 performance_insights_kms_key_id = aws_kms_key.pi_key.arn

 # Use optimized parameters
 parameter_group_name = aws_db_parameter_group.contact_db_params.name

 # Enable enhanced monitoring
 monitoring_interval = 60
 monitoring_role_arn = aws_iam_role.rds_monitoring.arn
}

```

## 5. Performance Monitoring

CloudWatch dashboard: [infra/modules/monitoring/main.tf](#) lines 130-180

```

resource "aws_cloudwatch_dashboard" "performance" {
 dashboard_name = "performance-metrics"

 dashboard_body = jsonencode({
 widgets = [
 {
 type = "metric"
 width = 12
 height = 6
 properties = {

```

```

 metrics = [
 ["AWS/ApiGateway", "Latency", "ApiName",
aws_api_gateway_rest_api.contact_api.name],
 ["AWS/ApiGateway", "CacheHitCount", "ApiName",
aws_api_gateway_rest_api.contact_api.name],
 ["AWS/ApiGateway", "CacheMissCount", "ApiName",
aws_api_gateway_rest_api.contact_api.name]
]
 period = 300
 stat = "Average"
 region = data.aws_region.current.name
 title = "API Gateway Performance"
 }
},
{
 type = "metric"
 width = 12
 height = 6
 properties = {
 metrics = [
 ["AWS/Lambda", "Duration", "FunctionName",
aws_lambda_function.contact_form.name],
 ["AWS/Lambda", "ConcurrentExecutions", "FunctionName",
aws_lambda_function.contact_form.name],
 ["AWS/Lambda", "ProvisionedConcurrencySpillover",
"FunctionName", aws_lambda_function.contact_form.name]
]
 period = 300
 stat = "Average"
 region = data.aws_region.current.name
 title = "Lambda Performance"
 }
}
]
})
}

```

## 5) Cost Optimization

### Current state

- CloudWatch billing alarm configured; artifacts bucket versioning enabled.
- RDS configured for free-tier friendly options; final snapshot skipped to reduce cost.

### Gaps

- S3 lifecycle rules for website/artifacts removed/commented; no asset TTL strategy.
- Always-on RDS; no CloudFront log analysis for cost vs value.

### TF improvements

- Re-enable S3 lifecycle for noncurrent versions and multipart cleanup; set appropriate Cache-Control for static assets.
- Add cost-focused dashboards/alarms (S3 bytes, Lambda duration, API Gateway costs); right-size resources regularly.
- Consider Aurora Serverless v2 if persistence needs grow with variable load.

## Evidence

- [infra/modules/monitoring/main.tf](#) (billing alarm)
- [infra/modules/s3/main.tf](#) (artifacts versioning; lifecycle commented out)
- [infra/modules/rds/main.tf](#) (free-tier settings)

## 6) Sustainability

### Current state

- Static hosting with CDN and serverless compute reduces idle waste; web images include WebP formats.
- Use S3 Intelligent-Tiering for storage to optimize energy and cost based on access patterns.
- Enable S3 lifecycle rules for all buckets to automatically delete old versions and unused objects.
- Monitor and report carbon footprint using AWS CloudWatch and third-party tools (e.g., AWS Customer Carbon Footprint Tool).

### Gaps

- RDS instance is always-on; caching/TTLs not optimized; no autoscaling or scheduled scale-down.
- Higher energy consumption and carbon footprint due to always-on resources and inefficient scaling.
- Missed opportunities for cost savings with S3 storage and lifecycle management.

### TF improvements

- Optimize CloudFront and S3 caching policies (set appropriate Cache-Control headers, enable Brotli/gzip compression).
- Add Terraform resources for monitoring and reporting carbon footprint (e.g., CloudWatch dashboards, custom metrics).
- Tag all resources with sustainability-related metadata (e.g., "environment", "purpose", "owner") for tracking and reporting.
- Audit and optimize static assets (images, JS, CSS) during CI/CD; automate with Terraform and build steps.

## Evidence

- [web/static/images/\\*.webp](#) (optimized images)
- [infra/modules/cloudfront/main.tf](#) (caching)

### Sustainability improvements (code refs)

## 1. CloudFront Compression & Caching

-Enabled Brotli/gzip compression and tiered TTLs for static assets.

```
default_cache_behavior {
 allowed_methods = ["GET", "HEAD", "OPTIONS"]
 cached_methods = ["GET", "HEAD", "OPTIONS"]
 target_origin_id = "s3-origin"
 compress = true

 forwarded_values {
 query_string = false
 cookies {
 forward = "none"
 }
 headers = ["Origin", "Access-Control-Request-Method", "Access-
Control-Request-Headers"]
 }

 viewer_protocol_policy = "redirect-to-https"
 min_ttl = 0
 default_ttl = 86400 # 24 hours
 max_ttl = 31536000 # 1 year
 response_headers_policy_id =
aws_cloudfront_response_headers_policy.optimized.id
}
```

## 2. S3 Lifecycle Rules

-Automatic cleanup of old versions and incomplete multipart uploads.

```
resource "aws_s3_bucket_lifecycle_configuration" "website_lifecycle" {
 bucket = aws_s3_bucket.website.id
 rule {
 id = "cleanup_old_versions"
 status = "Enabled"
 filter {
 prefix = ""
 }
 noncurrent_version_expiration {
 noncurrent_days = 30
 }
 abort_incomplete_multipart_upload {
 days_after_initiation = 7
 }
 }
}
```

### 3. Carbon Footprint Monitoring

-CloudWatch dashboard and custom metrics for sustainability reporting.

```
properties = {
 metrics = [
 ["AWS/Usage", "CarbonFootprint", "Service", "EC2", { stat =
"Sum" }],
 ["AWS/Usage", "CarbonFootprint", "Service", "S3", { stat =
"Sum" }],
 ["AWS/Usage", "CarbonFootprint", "Service", "Lambda", { stat
= "Sum" }]
]
 period = 86400
 region = data.aws_region.current.name
 title = "AWS Carbon Footprint (Daily)"
}
```

### 4. Resource Tagging

-Sustainability-related metadata tags (environment, purpose, owner) on all resources.

```
Common tags
locals {
 common_tags = {
 Environment = "development"
 Project = "assignment"
 ManagedBy = "terraform"
 Purpose = "sustainability"
 Owner = "DevOpsTeam"
 Sustainability = "true"
 }
}
```

### 5. CI/CD Asset Optimization

-Automated image, JS, and CSS optimization in build pipeline.

-buildspec-web.yml

```
Optimize static assets (images, JS, CSS)
- echo "Optimizing static assets..."
- npm --prefix web run optimize:assets || echo "Asset optimization
skipped (no script)"
- echo "Building static site..."
```

```
"optimize:assets": "npx imagemin static/images/* --out-dir=static/images && npx terser static/js/*.js -o static/js/ --compress --mangle && npx postcss static/css/*.css -o static/css/"
```

---

## Action Items (Optional)

- Priority P0:
  - Remove RDS public ingress and move DB to private subnets with strict SGs.
  - Add API authentication (e.g., Cognito/JWT) and attach AWS WAF to API.
  - Add CloudWatch alarms + SNS for CodeBuild/CodePipeline failures and Lambda errors.
- Priority P1:
  - Place Lambda in VPC and restrict egress; add DLQ and reserved concurrency.
  - Re-enable S3 lifecycle for website/artifacts; set Cache-Control headers via deploy step.
  - Tighten IAM policies for CodeBuild/CodePipeline to least privilege.
- Priority P2:
  - Enable API Gateway access logs/caching; add dashboards for API and Lambda.
  - Evaluate Multi-AZ for RDS and Performance Insights; consider Aurora Serverless v2.
  - Add Terratest-based infra tests in CI and manual approval gates for prod.

## Notes (Optional)

-