

# Well-Architected Framework (WAF) Review – DevOps CI/CD Project

---

## Instructions

1. Baseline: Briefly describe your existing setup (tools, repos, target environment, runtime). 15 sentences max.
  2. Per pillar: Fill the sections below:
    - Current state
    - Gaps
    - TF improvements
    - Evidence
  3. Terraform: Implement or reference TF changes that realize improvements.
  4. Validation: Link evidence (TF code lines, logs, screenshots) where possible.
- 

## Baseline

- Tools: AWS (S3, CloudFront, API Gateway, Lambda, RDS, CloudWatch), CodePipeline, CodeBuild, SSM Parameter Store, Terraform, Go (tests), Node.js 18 (web, lambda), ESLint, Stylelint, Vitest.
  - Repositories/Monorepo: Single repo with **infra/** (Terraform), **cicd/** (pipelines), and **web/** (static site + lambda).
  - Environments (dev/stage/prod): Single "development" environment tags; pipelines branch-driven (**develop**).
  - Cloud provider/region(s): AWS **us-east-1** (set in buildspecs and Terraform backend).
  - Runtimes (frontend/backend/lambda/DB): Static web (HTML/CSS/JS), Node.js Lambda for contact form, PostgreSQL RDS, API Gateway REST API.
  - CI/CD (build, test, deploy): Two CodePipelines (infra & web) with CodeBuild steps; linting and tests in **buildspec-\*.yaml**; S3 deploy, Lambda update, CloudFront invalidation.
  - IaC (Terraform versions/modules): Terraform **>=1.3** (backend S3/DynamoDB); modules for **s3**, **cloudfront**, **rds**, **lambda**, **api-gateway**, **iam**, **monitoring**, and **cicd**.
  - Observability (logs/metrics/tracing): CloudWatch logs for Lambda via **AWSLambdaBasicExecutionRole**; CloudWatch billing alarm; CodeBuild/CodePipeline logs.
  - Networking (VPCs/subnets): Uses default VPC only for RDS security group; Lambda not in VPC; API Gateway public; CloudFront over S3 OAI.
  - Security (IAM/KMS/secrets mgmt): IAM roles for CodeBuild/CodePipeline/Lambda; SSM parameters for DB creds; S3 public access blocked; CloudFront OAI for S3.
  - Data stores (RDS/S3/others): RDS Postgres instance; S3 website and artifacts buckets; SSM Parameter Store for config.
  - Edge/CDN: CloudFront distribution with HTTPS redirect; OAI restricting S3.
  - Key SLIs/SLOs: Not explicitly defined in repo.
- 

## 1) Operational Excellence

## Current state

- CI/CD via two CodePipelines with CodeBuild projects; infra pipeline runs fmt/validate/plan/apply; web pipeline runs lint/tests and deploys static site + Lambda + CF invalidation.
- Common tagging in Terraform locals; basic Go test scaffold present for infra.

## Gaps

- No alarms/notifications on pipeline or build failures; no manual approval gates.
- Limited automated tests (Terratest skipped); no runbooks.

## TF improvements

- Add CloudWatch alarms + SNS topics for CodeBuild/CodePipeline failures; add manual approval stage in pipelines.
- Add Terratest stage and enforce on PRs; expand tagging (owner, cost-center, service, environment).

## Evidence

- `cicd/main.tf` (CodePipeline, CodeBuild stages)
  - `buildspec-infra.yml`, `buildspec-web.yml` (lint/test/plan/apply/deploy)
  - `infra/main.tf` locals `common_tags`
  - `infra/tests/infra_integration_test.go`
- 

# 2) Security

## Current state

- S3 website bucket blocks public access; CloudFront OAI policy restricts reads.
- API Gateway POST/OPTIONS without auth; RDS SG allows 0.0.0.0/0 to 5432 (demo).
- Terraform remote state in S3 with DynamoDB lock and encryption.
- DB credentials stored in SSM Parameter Store (`/rds/db_username`, `/rds/db_password` SecureString, `/rds/db_name`, `/rds/rds_address`) and read by Lambda at runtime.
- There are no Network ACLs (NACLs) configured for your VPC subnets. All subnet-level traffic filtering is handled by default AWS settings and security groups.

## Gaps

- RDS publicly accessible SG; Lambda not in VPC; no KMS CMKs for S3/RDS.
- API lacks authentication/authorization and WAF; CodeBuild IAM policies broad with wildcards.
- No secret rotation; SSM path not scoped to environment.
- Missing stateless, subnet-level traffic filtering
- No defense-in-depth at the subnet layer
- Cannot explicitly deny unwanted IPs or ports at the subnet level

## TF improvements

- Place RDS in private subnets and restrict SG to Lambda/VPC CIDR; attach KMS CMK to RDS and S3.
- Put Lambda in VPC with least-priv SG; add Secrets Manager for credentials with rotation; scope SSM paths by env.
- Add API auth (API key/JWT/Cognito) and AWS WAF ACL; tighten IAM to least privilege.
- Define a NACL resource in Terraform

## Evidence

- `infra/modules/s3/main.tf` (public access block, OAI policy)
- `infra/modules/cloudfront/main.tf` (HTTPS redirect, OAI)
- `infra/modules/lambda/main.tf` (IAM role, SSM policy)
- `infra/modules/api-gateway/main.tf` (no auth)
- `infra/modules/rds/main.tf` (public SG 0.0.0.0/0)
- `infra/backend.tf` (S3 backend with DynamoDB lock)

## Improvements made (code refs)

S3 website bucket versioning: `infra/modules/s3/main.tf` lines 9-15

```
resource "aws_s3_bucket_versioning" "website_versioning" {
  bucket = aws_s3_bucket.website.id
  versioning_configuration {
    status = "Enabled"
  }
}
```

S3 website bucket encryption: `infra/modules/s3/main.tf` lines 17-26

```
resource "aws_s3_bucket_server_side_encryption_configuration"
"website_encryption" {
  bucket = aws_s3_bucket.website.id
  rule { apply_server_side_encryption_by_default { sse_algorithm =
"AES256" } }
}
```

S3 artifacts versioning: `infra/modules/s3/main.tf` lines 79-84

```
resource "aws_s3_bucket_versioning" "codepipeline_artifacts_versioning"
{
  bucket = aws_s3_bucket.codepipeline_artifacts.id
  versioning_configuration {
    status = "Enabled"
  }
}
```

S3 artifacts encryption: [infra/modules/s3/main.tf](#) lines 86-95

```
resource "aws_s3_bucket_server_side_encryption_configuration"
"codepipeline_artifacts_encryption" {
  bucket = aws_s3_bucket.codepipeline_artifacts.id

  rule {
    apply_server_side_encryption_by_default {
      sse_algorithm = "AES256"
    }
  }
}
```

CloudFront security headers policy: [infra/modules/cloudfront/main.tf](#) lines 20-37

```
default_cache_behavior {
  allowed_methods  = ["GET", "HEAD"]
  cached_methods  = ["GET", "HEAD"]
  target_origin_id = "s3-origin"

  forwarded_values {
    query_string = false
    cookies {
      forward = "none"
    }
  }
}

viewer_protocol_policy    = "redirect-to-https"
min_ttl                   = 0
default_ttl                = 3600
max_ttl                   = 86400
response_headers_policy_id = "60669652-455b-4ae9-85a4-c4c02393f86c" #
AWSManagedSecurityHeadersPolicy
}
```

CloudFront minimum TLS version: [infra/modules/cloudfront/main.tf](#) lines 45-48

```
viewer_certificate {
  cloudfront_default_certificate = true
  minimum_protocol_version       = "TLSv1.2_2021"
}
```

API Gateway access logs: [infra/modules/api-gateway/main.tf](#) lines 101-123

```

resource "aws_api_gateway_stage" "contact_stage" {
  deployment_id = aws_api_gateway_deployment.contact_deployment.id
  rest_api_id   = aws_api_gateway_rest_api.contact_api.id
  stage_name    = var.stage_name

  access_log_settings {
    destination_arn = aws_cloudwatch_log_group.api_gw_logs.arn
    format = jsonencode({
      requestId      = "$context.requestId",
      ip             = "$context.identity.sourceIp",
      caller         = "$context.identity.caller",
      user           = "$context.identity.user",
      requestTime    = "$context.requestTime",
      httpMethod     = "$context.httpMethod",
      resourcePath   = "$context.resourcePath",
      status         = "$context.status",
      protocol       = "$context.protocol",
      responseLength = "$context.responseLength",
      integrationStatus = "$context.integration.status",
      integrationError = "$context.integrationErrorMessage"
    })
  }
}

```

API Gateway log group and method throttling: [infra/modules/api-gateway/main.tf](#) lines 125-145

```

resource "aws_cloudwatch_log_group" "api_gw_logs" {
  name =
"/apigw/${aws_api_gateway_rest_api.contact_api.id}/${var.stage_name}"
  retention_in_days = 14
  tags = var.tags
}

resource "aws_api_gateway_method_settings" "all" {
  rest_api_id = aws_api_gateway_rest_api.contact_api.id
  stage_name  = aws_api_gateway_stage.contact_stage.stage_name
  method_path = "*/*"

  settings {
    metrics_enabled      = true
    logging_level        = "INFO"
    data_trace_enabled   = true
    throttling_burst_limit = 5
    throttling_rate_limit = 10
  }
}

```

```
resource "aws_wafv2_web_acl" "apigw_acl" {
  name          = "apigw-basic-acl"
  description   = "Basic protections for API Gateway"
  scope         = "REGIONAL"
  default_action {
    allow {}
  }
  rule {
    name          = "AWS-AWSManagedRulesCommonRuleSet"
    priority      = 1
    statement {
      managed_rule_group_statement {
        name          = "AWSManagedRulesCommonRuleSet"
        vendor_name   = "AWS"
      }
    }
    visibility_config {
      cloudwatch_metrics_enabled = true
      metric_name                 = "AWSManagedRulesCommonRuleSet"
      sampled_requests_enabled    = true
    }
  }
  visibility_config {
    cloudwatch_metrics_enabled = true
    metric_name                 = "apigw-acl"
    sampled_requests_enabled    = true
  }
  tags = var.tags
}

resource "aws_wafv2_web_acl_association" "apigw_acl_assoc" {
  resource_arn = aws_api_gateway_stage.contact_stage.arn
  web_acl_arn  = aws_wafv2_web_acl.apigw_acl.arn
}
```

```
resource "aws_iam_role_policy_attachment" "lambda_vpc_access" {
  role          = aws_iam_role.lambda_exec.name
  policy_arn    = "arn:aws:iam::aws:policy/service-role/AWSLambdaVPCAccessExecutionRole"
}
```

```

resource "aws_security_group" "lambda_sg" {
  name_prefix = "lambda-sg-"
  description = "Security group for Lambda to access RDS"
  vpc_id      = data.aws_vpc.default.id

  egress {
    from_port = 0
    to_port   = 0
    protocol  = "-1"
    cidr_blocks = ["0.0.0.0/0"]
  }

  tags = merge(var.tags, { Name = "lambda-sg" })
}

```

Lambda VPC config on function: [infra/modules/lambda/main.tf](#) lines 89-92

```

vpc_config {
  subnet_ids          = data.aws_subnets.default_vpc_subnets.ids
  security_group_ids = [aws_security_group.lambda_sg.id]
}

```

RDS SG ingress from Lambda SG: [infra/modules/rds/main.tf](#) lines 6-17

```

resource "aws_security_group" "rds_ingress" {
  name_prefix = "rds-ingress-5432-"
  description = "Allow inbound to Postgres from Lambda SG"
  vpc_id      = data.aws_vpc.default.id

  ingress {
    from_port = 5432
    to_port   = 5432
    protocol  = "tcp"
    security_groups = [var.allowed_sg_id]
    description    = "Allow Postgres from Lambda security
group"
  }
}

```

RDS private/encrypted + SG attach: [infra/modules/rds/main.tf](#) lines 48-58

```

resource "aws_db_instance" "contact_db" {
  storage_encrypted = var.storage_encrypted
  publicly_accessible = var.publicly_accessible
}

```

```
vpc_security_group_ids = [aws_security_group.rds_ingress.id]
}
```

IAM narrowed S3 bucket resources: [infra/modules/iam/main.tf](#) lines 276-279

```
Resource = [
    var.artifacts_bucket_arn,
    var.website_bucket_arn
]
```

Secrets Manager secret and version: [infra/main.tf](#) lines 58-75

```
resource "aws_secretsmanager_secret" "db_credentials" {
    name          = "project3/db-credentials"
    description    = "Database credentials for contact form"
    tags          = local.common_tags
}

resource "aws_secretsmanager_secret_version" "db_credentials_version" {
    secret_id      = aws_secretsmanager_secret.db_credentials.id
    secret_string  = jsonencode({
        username = coalesce(var.db_username,
data.aws_ssm_parameter.db_username.value)
        password = coalesce(var.db_password,
data.aws_ssm_parameter.db_password.value)
        host     = module.rds.rds_address
        database = coalesce(var.db_name,
data.aws_ssm_parameter.db_name.value)
        port     = module.rds.rds_port
    })
    depends_on    = [module.rds]
}
```

- Rotation function and rotation rule: [infra/main.tf](#) lines 76-103

```
resource "aws_serverlessapplicationrepository_cloudformation_stack"
"rds_rotation" {
    name          = "project3-rds-rotation"
    application_id = "arn:aws:serverlessrepo:us-east-
1:297356227824:applications/SecretsManagerRDSPostgreSQLRotationSingleUse
r"
    capabilities  = ["CAPABILITY_NAMED_IAM"]
    parameters = {
        functionName      = "project3-rds-rotation"
        vpcSubnetIds       = join(",",
data.aws_subnets.default_vpc_subnets.ids)
    }
}
```



```

    vpcSecurityGroupIds = module.lambda.lambda_security_group_id
  }
  semantic_version = "1.1.188"
  tags              = local.common_tags
}

resource "aws_secretsmanager_secret_rotation" "db_rotation" {
  secret_id          = aws_secretsmanager_secret.db_credentials.id
  rotation_lambda_arn =
aws_serverlessapplicationrepository_cloudformation_stack.rds_rotation.ou
tputs["RotationLambdaARN"]
  rotation_rules { automatically_after_days = 30 }
  depends_on =
[aws_secretsmanager_secret_version.db_credentials_version]
}

```

- Lambda environment + IAM for secret read: [infra/modules/lambda/main.tf](#) lines 56-66, 70-90

```

environment {
  variables = {
    ENVIRONMENT = "development"
    DB_SECRET_ARN = var.db_secret_arn
  }
}

```

```

resource "aws_iam_role_policy" "lambda_secrets_policy" {
  name = "lambda-secrets-access"
  role = aws_iam_role.lambda_exec.id
  policy = jsonencode({
    Version = "2012-10-17",
    Statement = [{
      Effect = "Allow",
      Action =
["secretsmanager:GetSecretValue","secretsmanager:DescribeSecret"],
      Resource = var.db_secret_arn
    }]
  })
}

```

- Lambda code reads from Secrets Manager with SSM fallback: [web/lambda/index.js](#) lines 1-18, 21-40, 44-74

```

// Security note: Database credentials are retrieved from AWS SSM
Parameter Store, not environment variables.
import { Client } from "pg";

```

```
import { SSMClient, GetParameterCommand } from "@aws-sdk/client-ssm";
import { SecretsManagerClient, GetSecretValueCommand } from "@aws-
sdk/client-secrets-manager";

// Initialize AWS clients
const region = process.env.AWS_REGION || 'us-east-1';
const ssmClient = new SSMClient({ region });
const secretsClient = new SecretsManagerClient({ region });
```

```
// Cache for database credentials to avoid repeated SSM calls
let dbCredentials = null;

// Function to get database credentials (Secrets Manager preferred,
// fallback to SSM)
async function getDbCredentials() {
  if (dbCredentials) {
    return dbCredentials;
  }

  try {
    const secretArn = process.env.DB_SECRET_ARN;
    if (secretArn) {
      const secretData = await secretsClient.send(new
GetSecretValueCommand({ SecretId: secretArn }));
      const secret = JSON.parse(secretData.SecretString || '{}');
      dbCredentials = { host: secret.host, user: secret.username,
password: secret.password, database: secret.database };
    } else {
      // Fallback to SSM parameters
```

```
const [dbHost, dbUser, dbPass, dbName] = await Promise.all([
  ssmClient.send(new GetParameterCommand({ Name:
"/rds/rds_address" })),
  ssmClient.send(new GetParameterCommand({ Name:
"/rds/db_username" })),
  ssmClient.send(new GetParameterCommand({ Name:
"/rds/db_password", WithDecryption: true })),
  ssmClient.send(new GetParameterCommand({ Name: "/rds/db_name"
}))
]);
dbCredentials = { host: dbHost.Parameter.Value, user:
dbUser.Parameter.Value, password: dbPass.Parameter.Value, database:
dbName.Parameter.Value };
return dbCredentials;
} catch (error) {
  console.error("Failed to retrieve database credentials:", error);
  throw new Error("Database configuration error");
}
```

```
}  
}
```

-Network ACLs (NACLs) have been implemented in your Terraform VPC module.[infra/vpc/main.tf](#)  
line 124-194

```
# Network ACLs for Public Subnets  
resource "aws_network_acl" "public" {  
  vpc_id = aws_vpc.main.id  
  tags = merge(var.tags, { Name = "${var.environment}-public-nacl" })  
}  
  
# Allow HTTPS inbound, deny all else (example)  
resource "aws_network_acl_rule" "public_https_inbound" {  
  network_acl_id = aws_network_acl.public.id  
  rule_number    = 100  
  egress         = false  
  protocol       = "tcp"  
  rule_action    = "allow"  
  cidr_block     = "0.0.0.0/0"  
  from_port     = 443  
  to_port        = 443  
}  
  
resource "aws_network_acl_rule" "public_deny_all_inbound" {  
  network_acl_id = aws_network_acl.public.id  
  rule_number    = 200  
  egress         = false  
  protocol       = "-1"  
  rule_action    = "deny"  
  cidr_block     = "0.0.0.0/0"  
  from_port     = 0  
  to_port        = 0  
}  
  
# Associate Public NACL with Public Subnets  
resource "aws_network_acl_association" "public" {  
  count          = length(var.public_subnet_cidrs)  
  subnet_id     = aws_subnet.public[count.index].id  
  network_acl_id = aws_network_acl.public.id  
}  
  
# Network ACLs for Private Subnets  
resource "aws_network_acl" "private" {  
  vpc_id = aws_vpc.main.id  
  tags = merge(var.tags, { Name = "${var.environment}-private-nacl" })  
}  
  
# Allow DB traffic from Lambda SG CIDR (example: adjust as needed)  
resource "aws_network_acl_rule" "private_db_inbound" {
```

```

    network_acl_id = aws_network_acl.private.id
    rule_number    = 100
    egress         = false
    protocol       = "tcp"
    rule_action    = "allow"
    cidr_block     = "10.0.0.0/8" # Example CIDR, adjust to Lambda SG
  subnet
    from_port     = 5432
    to_port       = 5432
  }

  resource "aws_network_acl_rule" "private_deny_all_inbound" {
    network_acl_id = aws_network_acl.private.id
    rule_number    = 200
    egress         = false
    protocol       = "-1"
    rule_action    = "deny"
    cidr_block     = "0.0.0.0/0"
    from_port      = 0
    to_port        = 0
  }

  # Associate Private NACL with Private Subnets
  resource "aws_network_acl_association" "private" {
    count          = length(var.private_subnet_cidrs)
    subnet_id      = aws_subnet.private[count.index].id
    network_acl_id = aws_network_acl.private.id
  }

```

### 3) Reliability

#### Current state

- Single RDS instance in primary region without Multi-AZ
- Basic RDS backups and maintenance windows configured
- No cross-region disaster recovery setup
- No automated failover mechanism
- RPO/RTO requirements not met (e-commerce needs: RPO 1h, RTO 4h)

#### Gaps

- Single point of failure with non-Multi-AZ RDS
- No cross-region redundancy for disaster recovery
- Missing health checks and automated failover
- No warm standby setup to meet RTO requirement
- Backup strategy insufficient for RPO requirement
- No DLQ or retries on Lambda functions
- Missing monitoring and alerting system

## TF improvements

### 1. Implement Warm Standby Architecture:

- Deploy standby RDS in us-west-2
- Configure cross-region replication
- Set up Route53 health checks and failover routing
- Enable Multi-AZ for primary RDS

### 2. Enhance Monitoring and Recovery:

- Add CloudWatch alarms for API errors (4XX/5XX)
- Monitor Lambda performance and failures
- Create operational dashboards
- Implement automated failover testing

### 3. Improve Data Protection:

- Configure RDS automated backups every 15 minutes
- Enable point-in-time recovery
- Implement cross-region S3 replication
- Set up proper backup retention policies

## Evidence

- `infra/modules/monitoring/main.tf` (only billing alarm)
- `infra/modules/api-gateway/main.tf` (stage definition)
- `infra/modules/lambda/main.tf` (no DLQ)
- `infra/modules/rds/main.tf` (backup/deletion flags)

## Reliability improvements (code refs)

### Warm Standby Architecture Decision

The decision to implement a warm standby architecture was based on the following requirements and considerations:

#### 1. Recovery Time Objective (RTO):

- Requirement: 4 hours maximum downtime allowed
- Warm standby provides faster recovery compared to cold standby or backup/restore
- Pre-provisioned infrastructure reduces deployment time during failover

#### 2. Recovery Point Objective (RPO):

- Requirement: Maximum 1 hour of data loss acceptable
- Continuous replication of data to standby region meets this requirement
- S3 cross-region replication for static content
- Database replication for dynamic data

### 3. Traffic Pattern Analysis:

- Steady, predictable traffic pattern
- Non-spiky workload suits warm standby's cost-effectiveness
- Lower cost compared to active-active while meeting RPO/RTO

### 4. Cost-Benefit Analysis:

- Warm standby provides optimal balance between recovery speed and cost
- Standby resources can run on smaller instances to reduce costs
- No need for complex active-active synchronization

### 5. Operational Complexity:

- Simpler than active-active architecture
- Automated failover through Route53 health checks
- Clear, well-defined failover process
- Easier to test and maintain

## Summary of Reliability Improvements

We implemented a warm standby architecture and enhanced reliability with the following changes:

- **Created standby RDS module in a different region:**
  - Folder: `infra/modules/rds-standby/`
  - Files: `main.tf`, `variables.tf`, `outputs.tf`
  - Purpose: Deploys a scaled-down standby PostgreSQL RDS instance in `us-west-2`.
- **Enabled S3 cross-region replication for static assets:**
  - File: `infra/modules/s3/replication.tf`
  - Purpose: Replicates website bucket to standby region for DR.
- **Configured Route53 DNS failover for APIs:**
  - File: `infra/modules/route53/failover.tf`
  - Supporting files: `variables.tf`, `outputs.tf`
  - Purpose: Automated DNS failover between primary and standby API endpoints.
- **Documented the warm standby setup and failover process:**
  - File: `WARM-STANDBY-README.md` (project root)
  - Purpose: Instructions and validation steps for disaster recovery.

All new modules follow Terraform best practices with variables, outputs, and clear separation of primary/standby resources.

Key code implementations:

#### 1. RDS Standby Configuration:

```
# infra/modules/rds-standby/main.tf
resource "aws_db_instance" "contact_db_standby" {
  identifier          = "contact-db-standby"
  engine              = "postgres"
  instance_class      = "db.t3.micro" # Scaled down for cost in standby
  allocated_storage   = 20
  storage_encrypted   = true

  backup_retention_period = 7
  backup_window           = "03:00-04:00"
  maintenance_window      = "Mon:04:00-Mon:05:00"

  multi_az            = false # Single AZ for standby to reduce costs
  publicly_accessible = false

  vpc_security_group_ids = [aws_security_group.rds_standby_sg.id]
  db_subnet_group_name   = aws_db_subnet_group.standby_subnet_group.name

  tags = merge(var.tags, {
    Name = "contact-db-standby"
    Role = "warm-standby"
  })
}
```

## 2. S3 Cross-Region Replication:

```
# infra/modules/s3/replication.tf
resource "aws_s3_bucket_replication_configuration" "website_replication"
{
  role  = aws_iam_role.replication_role.arn
  bucket = aws_s3_bucket.website.id

  rule {
    id      = "website-standby-replication"
    status  = "Enabled"

    destination {
      bucket      = aws_s3_bucket.website_standby.arn
      storage_class = "STANDARD"
    }
  }
}
```

## 3. Route53 DNS Failover:

```
# infra/modules/route53/failover.tf
resource "aws_route53_health_check" "primary_api" {
  fqdn      = var.primary_api_domain
```

```

port          = 443
type          = "HTTPS"
resource_path  = "/health"
failure_threshold = "3"
request_interval = "30"

tags = merge(var.tags, {
  Name = "primary-api-health-check"
})
}

resource "aws_route53_record" "api_primary" {
  zone_id = var.hosted_zone_id
  name     = var.api_domain
  type     = "A"

  failover_routing_policy {
    type = "PRIMARY"
  }

  alias {
    name           = var.primary_api_domain
    zone_id        = var.primary_api_zone_id
    evaluate_target_health = true
  }

  health_check_id = aws_route53_health_check.primary_api.id
  set_identifier   = "primary"
}

resource "aws_route53_record" "api_secondary" {
  zone_id = var.hosted_zone_id
  name     = var.api_domain
  type     = "A"

  failover_routing_policy {
    type = "SECONDARY"
  }

  alias {
    name           = var.standby_api_domain
    zone_id        = var.standby_api_zone_id
    evaluate_target_health = true
  }

  set_identifier = "secondary"
}

---

- Lambda reserved concurrency: `infra/modules/lambda/main.tf` lines 109-
111
```109:111:infra/modules/lambda/main.tf

```



```
# Reserve concurrency to prevent thundering herds and protect DB
reserved_concurrent_executions = 5
```

- Lambda DLQ and invoke config: [infra/modules/lambda/main.tf](#) lines 130-145

```
resource "aws_sqs_queue" "lambda_dlq" {
  name = "contact-form-dlq"
  tags = var.tags
}

resource "aws_lambda_event_invoke_config" "contact_eic" {
  function_name = aws_lambda_function.contact.function_name
  destination_config { on_failure { destination =
aws_sqs_queue.lambda_dlq.arn } }
  maximum_retry_attempts      = 2
  maximum_event_age_in_seconds = 3600
}
```

- Lambda errors alarm: [infra/modules/monitoring/main.tf](#) lines 25-42

```
resource "aws_cloudwatch_metric_alarm" "lambda_errors" {
  alarm_name          = "lambda-contact-errors"
  comparison_operator = "GreaterThanThreshold"
  evaluation_periods  = 1
  metric_name         = "Errors"
  namespace           = "AWS/Lambda"
  period              = 60
  statistic            = "Sum"
  threshold            = 1
  alarm_actions        = var.alarm_actions
  dimensions = { FunctionName = var.lambda_function_name }
}
```

- Lambda p95 duration alarm: [infra/modules/monitoring/main.tf](#) lines 44-60

```
resource "aws_cloudwatch_metric_alarm" "lambda_duration" {
  alarm_name          = "lambda-contact-duration-p95"
  comparison_operator = "GreaterThanThreshold"
  evaluation_periods  = 1
  metric_name         = "Duration"
  namespace           = "AWS/Lambda"
  period              = 60
  statistic            = "p95"
  threshold            = 3000
  alarm_actions        = var.alarm_actions
}
```

```
    dimensions = { FunctionName = var.lambda_function_name }
}
```

- API Gateway 5XX alarm: [infra/modules/monitoring/main.tf](#) lines 62-79

```
resource "aws_cloudwatch_metric_alarm" "apigw_5xx" {
  alarm_name          = "apigw-5xx-errors"
  comparison_operator = "GreaterThanThreshold"
  evaluation_periods  = 1
  metric_name         = "5XXError"
  namespace           = "AWS/ApiGateway"
  period              = 60
  statistic            = "Sum"
  threshold            = 1
  alarm_actions        = var.alarm_actions
  dimensions = { ApiId = var.api_gateway_id, Stage =
var.api_gateway_stage, Resource = "/contact", Method = "POST" }
}
```

- Multi-AZ for primary RDS and DMS cross-region replication:

```
resource "aws_db_instance" "contact_db" {
  # ...existing config...
  multi_az = true
  # Cross-region replication with AWS DMS
  # See DMS resources below
}
```

- DMS replication resources:

```
resource "aws_dms_replication_instance" "rds_replication" { ... }
resource "aws_dms_endpoint" "source" { ... }
resource "aws_dms_endpoint" "target" { ... }
resource "aws_dms_replication_task" "rds_to_standby" { ... }
resource "aws_dms_replication_subnet_group" "dms_subnet_group" { ... }
```

- DMS table mappings and task settings:

```
{
  "rules": [
    {
      "rule-type": "selection",
      "rule-id": "1",
      "rule-name": "1",
```

```

    "object-locator": {
      "schema-name": "%",
      "table-name": "%"
    },
    "rule-action": "include"
  }
]
}

```

```

{
  "TargetMetadata": { ... },
  "FullLoadSettings": { ... },
  "Logging": { ... },
  "ControlTablesSettings": { ... },
  "StreamBufferSettings": { ... },
  "ChangeProcessingDdlHandlingPolicy": { ... },
  "ErrorBehavior": { ... },
  "ChangeProcessingPolicy": { ... }
}

```

- S3 lifecycle for backup retention:

```

resource "aws_s3_bucket_lifecycle_configuration" "website_lifecycle" {
  bucket = aws_s3_bucket.website.id
  rule {
    id      = "cleanup_old_versions"
    status  = "Enabled"
    filter {
      prefix = ""
    }
    noncurrent_version_expiration {
      noncurrent_days = 30
    }
  }
}

```

---

## 4) Performance Efficiency

### Current state

Initial performance analysis revealed several areas needing optimization:

1. Content Delivery Network (CDN):
  - S3 + CloudFront distribution with basic 24-hour TTL
  - No compression enabled for static assets

- Missing Cache-Control headers for browser caching
- Default CloudFront settings without optimization

## 2. API Gateway and Lambda:

- Default Lambda configuration (128MB memory, 3s timeout)
- No API Gateway caching implemented
- Missing throttling and concurrency controls
- Basic database connection handling

## 3. Database Layer:

- RDS PostgreSQL on db.t3.micro (baseline tier)
- Default parameter group without tuning
- No Performance Insights monitoring
- Basic connection management

## 4. Performance Monitoring:

- Limited visibility into system performance
- No comprehensive monitoring dashboard
- Missing performance metrics collection
- No automated performance alerting

# Improvements Made

## 1. Content Delivery Network Optimization:

- Enabled CloudFront compression with Brotli support
- Implemented tiered TTL strategy:
  - Static assets: 1-year cache (31536000s)
  - Dynamic content: 24-hour cache (86400s)
  - API responses: 5-minute cache (300s)
- Added Cache-Control headers via CloudFront policies
- Configured optimal compression settings for asset types

## 2. API and Lambda Performance Enhancements:

- Increased Lambda memory to 256MB for better performance
- Extended timeout to 10s for reliable operation
- Added API Gateway caching (0.5GB cache size)
- Implemented database connection pooling
- Set up provisioned concurrency (2 instances)

## 3. Database Optimizations:

- Upgraded to db.t3.small for improved performance
- Enabled Performance Insights (7-day retention)
- Optimized database parameters:
  - Shared buffers: 20% of instance memory

- Work memory: 8MB per connection
- Max connections: 100
- Implemented enhanced monitoring (60s intervals)

#### 4. Performance Monitoring Framework:

- Created comprehensive CloudWatch dashboard
- Set up performance metric alarms
- Added API Gateway metrics tracking
- Implemented Lambda performance monitoring

### Evidence

- CloudFront basic config: [infra/modules/cloudfront/main.tf](#) lines 15-25:

```
default_cache_behavior {
  min_ttl = 0
  default_ttl = 86400 # 24 hours default
  max_ttl = 31536000 # 1 year max
}
```

- Lambda basic settings: [infra/modules/lambda/main.tf](#) lines 45-48:

```
resource "aws_lambda_function" "contact_form" {
  memory_size = 128 # Default memory
  timeout     = 3   # Default timeout
}
```

- API Gateway without caching: [infra/modules/api-gateway/main.tf](#) lines 89-95:

```
resource "aws_api_gateway_stage" "api" {
  # No cache_cluster_enabled setting
  # No method_settings for caching
}
```

- RDS base configuration: [infra/modules/rds/main.tf](#) lines 25-30:

```
resource "aws_db_instance" "contact_db" {
  instance_class = "db.t3.micro"
  # No performance_insights_enabled
  # Default parameter group
}
```

## 1. CloudFront Optimizations

CloudFront compression and caching: [infra/modules/cloudfront/main.tf](#) lines 30-52

```
resource "aws_cloudfront_distribution" "website" {
  # ... existing configuration ...

  default_cache_behavior {
    compress                = true
    viewer_protocol_policy = "redirect-to-https"

    min_ttl                = 0
    default_ttl            = 86400    # 24 hours
    max_ttl                = 31536000 # 1 year

    cached_methods        = ["GET", "HEAD", "OPTIONS"]
    allowed_methods        = ["GET", "HEAD", "OPTIONS"]

    forwarded_values {
      query_string = false
      cookies { forward = "none" }
      headers      = ["Origin", "Access-Control-Request-Method",
"Access-Control-Request-Headers"]
    }
  }

  ordered_cache_behavior {
    path_pattern      = "/assets/*"
    allowed_methods   = ["GET", "HEAD"]
    cached_methods    = ["GET", "HEAD"]
    target_origin_id  = aws_s3_bucket.website.id

    compress                = true
    viewer_protocol_policy = "redirect-to-https"

    min_ttl                = 86400    # 1 day
    default_ttl            = 604800    # 1 week
    max_ttl                = 31536000 # 1 year
  }
}
```

CloudFront response headers policy: [infra/modules/cloudfront/main.tf](#) lines 54-75

```
resource "aws_cloudfront_response_headers_policy" "security_headers" {
  name = "performance-security-headers"

  custom_headers_config {
```

```

    items {
      header    = "Cache-Control"
      value     = "public, max-age=31536000"
      override  = true
    }
    items {
      header    = "Accept-Encoding"
      value     = "gzip, deflate, br"
      override  = true
    }
  }

  security_headers_config {
    content_type_options {
      override = true
    }
    frame_options {
      frame_option = "DENY"
      override     = true
    }
    strict_transport_security {
      access_control_max_age_sec = 31536000
      include_subdomains         = true
      override                   = true
    }
  }
}

```

## 2. API Gateway Caching

API Gateway cache settings: [infra/modules/api-gateway/main.tf](#) lines 150-180

```

resource "aws_api_gateway_stage" "contact_stage" {
  # ... existing configuration ...

  cache_cluster_enabled = true
  cache_cluster_size    = "0.5" # 0.5GB cache

  variables = {
    "cacheEnabled" = "true"
  }
}

resource "aws_api_gateway_method_settings" "contact_cache" {
  rest_api_id = aws_api_gateway_rest_api.contact_api.id
  stage_name  = aws_api_gateway_stage.contact_stage.stage_name
  method_path = "*/*"

  settings {
    metrics_enabled = true
  }
}

```

```

    logging_level      = "INFO"
    caching_enabled    = true
    cache_ttl_in_seconds = 300 # 5 minutes cache

    throttling_burst_limit = 100
    throttling_rate_limit  = 50
  }
}

```

### 3. Lambda Optimization

Lambda performance configuration: [infra/modules/lambda/main.tf](#) lines 80-110

```

resource "aws_lambda_function" "contact_form" {
  # ... existing configuration ...

  memory_size = 256
  timeout     = 10

  environment {
    variables = {
      NODE_OPTIONS = "--enable-source-maps"
      POSTGRES_MAX_CONNECTIONS = "10"
    }
  }

  provisioned_concurrent_executions = 2

  vpc_config {
    subnet_ids          = data.aws_subnets.default_vpc_subnets.ids
    security_group_ids = [aws_security_group.lambda_sg.id]
  }
}

# Add connection pooling utility
resource "aws_lambda_layer_version" "pg_pool" {
  layer_name      = "pg-pool-layer"
  description     = "PostgreSQL connection pooling layer"
  filename        = "layers/pg-pool.zip"
  compatible_runtimes = ["nodejs18.x"]
}

```

### 4. RDS Performance

RDS performance configuration: [infra/modules/rds/main.tf](#) lines 85-120

```

resource "aws_db_parameter_group" "contact_db_params" {
  name = "contact-db-params"
}

```



```

family = "postgres14"

parameter {
  name  = "shared_buffers"
  value = "{DBInstanceClassMemory*20/100}" # 20% of instance memory
}

parameter {
  name  = "max_connections"
  value = "100"
}

parameter {
  name  = "work_mem"
  value = "8388608" # 8MB
}
}

resource "aws_db_instance" "contact_db" {
  # ... existing configuration ...

  instance_class = "db.t3.small"

  # Enable performance insights
  performance_insights_enabled      = true
  performance_insights_retention_period = 7
  performance_insights_kms_key_id   = aws_kms_key.pi_key.arn

  # Use optimized parameters
  parameter_group_name = aws_db_parameter_group.contact_db_params.name

  # Enable enhanced monitoring
  monitoring_interval = 60
  monitoring_role_arn = aws_iam_role.rds_monitoring.arn
}

```

## 5. Performance Monitoring

CloudWatch dashboard: [infra/modules/monitoring/main.tf](#) lines 130-180

```

resource "aws_cloudwatch_dashboard" "performance" {
  dashboard_name = "performance-metrics"

  dashboard_body = jsonencode({
    widgets = [
      {
        type  = "metric"
        width = 12
        height = 6
        properties = {

```

```

        metrics = [
            ["AWS/ApiGateway", "Latency", "ApiName",
aws_api_gateway_rest_api.contact_api.name],
            ["AWS/ApiGateway", "CacheHitCount", "ApiName",
aws_api_gateway_rest_api.contact_api.name],
            ["AWS/ApiGateway", "CacheMissCount", "ApiName",
aws_api_gateway_rest_api.contact_api.name]
        ]
        period = 300
        stat = "Average"
        region = data.aws_region.current.name
        title = "API Gateway Performance"
    }
},
{
    type = "metric"
    width = 12
    height = 6
    properties = {
        metrics = [
            ["AWS/Lambda", "Duration", "FunctionName",
aws_lambda_function.contact_form.name],
            ["AWS/Lambda", "ConcurrentExecutions", "FunctionName",
aws_lambda_function.contact_form.name],
            ["AWS/Lambda", "ProvisionedConcurrencySpillover",
"FunctionName", aws_lambda_function.contact_form.name]
        ]
        period = 300
        stat = "Average"
        region = data.aws_region.current.name
        title = "Lambda Performance"
    }
}
]
})
}

```

## 5) Cost Optimization

### Current state

- CloudWatch billing alarm configured; artifacts bucket versioning enabled.
- RDS configured for free-tier friendly options; final snapshot skipped to reduce cost.

### Gaps

- S3 lifecycle rules for website/artifacts removed/commented; no asset TTL strategy.
- Always-on RDS; no CloudFront log analysis for cost vs value.
- S3 cross-region replication uses STANDARD storage class instead of cost-effective options (IA/Glacier)

- No CloudFront price class optimization - using default (all edge locations globally)
- No automated resource scheduling (e.g., RDS stop/start based on usage patterns)
- Missing cost allocation tags for detailed cost tracking by environment/feature
- CloudWatch log retention set to 14 days but could be optimized based on compliance needs
- DMS replication instance running continuously (dms.t3.medium) for cross-region replication adds significant cost

## TF improvements

- Re-enable S3 lifecycle for noncurrent versions and multipart cleanup; set appropriate Cache-Control for static assets.
- Add cost-focused dashboards/alarms (S3 bytes, Lambda duration, API Gateway costs); right-size resources regularly.
- Consider Aurora Serverless v2 if persistence needs grow with variable load.

## Evidence

- `infra/modules/monitoring/main.tf` (billing alarm)
- `infra/modules/s3/main.tf` (artifacts versioning; lifecycle commented out)
- `infra/modules/rds/main.tf` (free-tier settings)

## Cost Optimization Improvements

### 1. S3 Intelligent-Tiering & Enhanced Lifecycle Rules

File: `infra/modules/s3/main.tf` - Lines 97-170

#### S3 Intelligent-Tiering Configuration

- Automatic cost optimization based on access patterns
- Archive tier after 90 days, Deep Archive after 180 days

```
# S3 Intelligent-Tiering configuration for automatic cost optimization
resource "aws_s3_bucket_intelligent_tiering_configuration"
"website_tiering" {
  bucket = aws_s3_bucket.website.id
  name    = "website-intelligent-tiering"

  tiering {
    access_tier = "ARCHIVE_ACCESS"
    days        = 90
  }

  tiering {
    access_tier = "DEEP_ARCHIVE_ACCESS"
    days        = 180
  }
}
```

## Enhanced Website Lifecycle Rules

- Multiple storage class transitions for optimal cost savings
- Automated cleanup of old versions and incomplete uploads

```
resource "aws_s3_bucket_lifecycle_configuration" "website_lifecycle" {
  bucket = aws_s3_bucket.website.id

  rule {
    id      = "cleanup_old_versions"
    status  = "Enabled"
    filter {
      prefix = ""
    }
    noncurrent_version_expiration {
      noncurrent_days = 30
    }
    abort_incomplete_multipart_upload {
      days_after_initiation = 7
    }
  }

  rule {
    id      = "transition_to_ia"
    status  = "Enabled"
    filter {
      prefix = ""
    }
    transition {
      days          = 30
      storage_class = "STANDARD_IA"
    }
    transition {
      days          = 90
      storage_class = "GLACIER"
    }
    transition {
      days          = 365
      storage_class = "DEEP_ARCHIVE"
    }
  }
}
```

## CI/CD Artifacts Lifecycle Optimization

- Aggressive cleanup for temporary CI/CD artifacts
- 90-day expiration for cost control

```
resource "aws_s3_bucket_lifecycle_configuration" "artifacts_lifecycle" {
  bucket = aws_s3_bucket.codepipeline_artifacts.id

  rule {
    id      = "cleanup_artifacts"
    status  = "Enabled"
    filter {
      prefix = ""
    }
    noncurrent_version_expiration {
      noncurrent_days = 7 # Shorter retention for CI/CD artifacts
    }
    abort_incomplete_multipart_upload {
      days_after_initiation = 1
    }
    expiration {
      days = 90 # Delete old artifacts after 90 days
    }
  }
}
```

---

## 2. Cross-Region Replication Storage Optimization

File: [infra/modules/s3/replication.tf](#) - Line 18

- Changed from STANDARD to STANDARD\_IA for 50% cost reduction on standby storage

```
destination {
  bucket      = aws_s3_bucket.website_standby.arn
  storage_class = "STANDARD_IA" # Cost optimization: Use IA for standby
  region
}
```

---

## 3. CloudFront Cost Optimization

File: [infra/modules/cloudfront/main.tf](#) - Lines 1-20

### CloudFront Access Logs S3 Bucket

- Dedicated bucket for CloudFront logs with automatic cleanup

```
# S3 bucket for CloudFront access logs
resource "aws_s3_bucket" "cloudfront_logs" {
  bucket      = "${var.s3_bucket_name}-cf-logs"
  force_destroy = true
}
```

```

tags          = var.tags
}

resource "aws_s3_bucket_lifecycle_configuration"
"cloudfront_logs_lifecycle" {
  bucket = aws_s3_bucket.cloudfront_logs.id

  rule {
    id      = "delete_old_logs"
    status  = "Enabled"
    filter {
      prefix = "cloudfront-logs/"
    }
    expiration {
      days = var.log_retention_days
    }
  }
}

```

File: [infra/modules/cloudfront/main.tf](#) - Lines 120-125

### Price Class & Access Logging Configuration

- Environment-based price class optimization
- CloudFront access logging for cost analysis

```

# Cost optimization: Use price class that covers US, Canada, Europe, and
Asia
price_class = var.price_class

# Enable access logging for cost analysis
logging_config {
  include_cookies = false
  bucket          = aws_s3_bucket.cloudfront_logs.bucket_domain_name
  prefix          = "cloudfront-logs/"
}

```

File: [infra/modules/cloudfront/variables.tf](#) - Lines 6-25

### Cost Optimization Variables

- Price class defaults to cost-effective PriceClass\_100
- Configurable log retention for compliance needs

```

variable "s3_bucket_name" {
  description = "Name of the S3 bucket for generating log bucket name"
  type        = string
}

```

```

}

variable "price_class" {
  description = "CloudFront price class for cost optimization"
  type        = string
  default     = "PriceClass_100" # US, Canada, Europe only
}

variable "log_retention_days" {
  description = "Number of days to retain CloudFront logs"
  type        = number
  default     = 30
}

```

---

## 4. Environment-Based Multi-AZ & RDS Optimization

File: [infra/variables.tf](#) - Lines 1-5

### Environment Variable for Cost Control

- Single variable controls cost optimizations across all resources

```

variable "environment" {
  description = "Environment name (development, staging, production)"
  type        = string
  default     = "development"
}

```

File: [infra/modules/rds/main.tf](#) - Lines 153-155

### Conditional Multi-AZ Configuration

- Multi-AZ only enabled in production for 50% cost savings in development

```

# Reliability improvements - conditional Multi-AZ based on environment
backup_retention_period = var.environment == "production" ? 7 : 1
multi_az                 = var.environment == "production" ? true : false

```

---

## 5. DMS Replication Instance Optimization

File: [infra/modules/rds/main.tf](#) - Lines 7-17

### Conditional DMS Instance

- DMS only created for production environment
- Right-sized instance classes and storage allocation

```
resource "aws_dms_replication_instance" "rds_replication" {
  count = var.environment == "production" ? 1 : 0 # Only create DMS for
production

  replication_instance_id      = "rds-replication-instance"
  allocated_storage            = var.environment == "production" ? 100 :
50 # Smaller storage for non-prod
  replication_instance_class   = var.environment == "production" ?
"dms.t3.medium" : "dms.t3.small"
  engine_version               = "3.4.6"
  publicly_accessible          = false
  multi_az                     = var.environment == "production" ? true :
false # Cost optimization
  vpc_security_group_ids      = [aws_security_group.rds_ingress.id]
  replication_subnet_group_id = var.dms_subnet_group_id
  tags = var.tags
}
```

File: [infra/modules/rds/main.tf](#) - Lines 21-53

### Conditional DMS Endpoints & Tasks

- All DMS components made conditional on production environment

```
resource "aws_dms_endpoint" "source" {
  count = var.environment == "production" ? 1 : 0 # Only create for
production

  endpoint_id    = "source-endpoint"
  endpoint_type  = "source"
  engine_name    = "postgres"
  username       = var.db_username
  password       = var.db_password
  server_name    = aws_db_instance.contact_db.address
  port          = 5432
  database_name  = var.db_name
  ssl_mode       = "require"
}

resource "aws_dms_endpoint" "target" {
  count = var.environment == "production" ? 1 : 0 # Only create for
production

  endpoint_id    = "target-endpoint"
  endpoint_type  = "target"
  engine_name    = "postgres"
}
```



```

username      = var.db_username
password      = var.db_password
server_name   = var.standby_rds_address
port          = 5432
database_name = var.db_name
ssl_mode      = "require"
}

resource "aws_dms_replication_task" "rds_to_standby" {
  count = var.environment == "production" ? 1 : 0 # Only create for
production

  replication_task_id      = "rds-to-standby"
  migration_type           = "cdc"
  replication_instance_arn =
aws_dms_replication_instance.rds_replication[0].arn
  source_endpoint_arn      = aws_dms_endpoint.source[0].arn
  target_endpoint_arn      = aws_dms_endpoint.target[0].arn
  table_mappings            = file("${path.module}/dms-table-
mappings.json")
  replication_task_settings = file("${path.module}/dms-task-
settings.json")
  tags = var.tags
}

```

---

## 6. Enhanced Cost Allocation Tags

File: `infra/main.tf` - Lines 18-35

### Comprehensive Tagging Strategy

- Cost allocation tags for detailed tracking by environment, feature, and owner
- Operational tags for automation and compliance

```

# Enhanced cost allocation tags
locals {
  common_tags = {
    # Cost allocation tags
    Environment      = var.environment
    Project          = "contact-form-webapp"
    ManagedBy        = "terraform"
    CostCenter        = "development"
    Owner            = "devops-team"
    BusinessUnit      = "engineering"
    Application       = "contact-form"

    # Operational tags
    Purpose          = "web-application"
    Sustainability   = "enabled"
  }
}

```

```

    AutoShutdown      = var.environment != "production" ? "enabled" :
"disabled"
    BackupRequired    = var.environment == "production" ? "yes" : "no"

    # Compliance tags
    DataClass         = "internal"
    Compliance        = "standard"
  }
}

```

---

## 7. CloudWatch Log Retention Optimization

File: `infra/modules/api-gateway/main.tf` - Line 163

### Configurable Log Retention

- Environment-based log retention periods for cost optimization

```

resource "aws_cloudwatch_log_group" "api_gw_logs" {
  name           =
"/apigw/${aws_api_gateway_rest_api.contact_api.id}/${var.stage_name}"
  retention_in_days = var.log_retention_days # Environment-based
retention
  tags           = var.tags
}

```

File: `infra/main.tf` - Lines 151-152

### Environment-Based Retention Configuration

- Production: 90 days, Development: 7 days for cost savings

```

module "api_gateway" {
  source = "../modules/api-gateway"

  api_name      = "contact-api"
  stage_name    = "dev"
  lambda_invoke_arn = module.lambda.lambda_invoke_arn
  aws_region    = var.aws_region
  log_retention_days = var.environment == "production" ? 90 : 7 #
Environment-based retention
  tags          = local.common_tags
}

```

---

## 8. Service-Specific Cost Monitoring

File: `infra/cost-optimization.tf` - Lines 1-110

### S3 Cost Alarm

- Environment-specific thresholds for cost control

```
resource "aws_cloudwatch_metric_alarm" "s3_costs" {
  alarm_name          = "s3-monthly-costs-${var.environment}"
  comparison_operator = "GreaterThanThreshold"
  evaluation_periods  = "1"
  metric_name         = "EstimatedCharges"
  namespace           = "AWS/Billing"
  period              = "86400"
  statistic            = "Maximum"
  threshold            = var.environment == "production" ? "50" : "10"
  alarm_description   = "S3 monthly costs exceeded threshold"
  alarm_actions       = []

  dimensions = {
    Currency    = "USD"
    ServiceName = "AmazonS3"
  }

  tags = local.common_tags
}
```

### Lambda Cost Alarm

- Proactive monitoring of Lambda execution costs

```
resource "aws_cloudwatch_metric_alarm" "lambda_costs" {
  alarm_name          = "lambda-monthly-costs-${var.environment}"
  comparison_operator = "GreaterThanThreshold"
  evaluation_periods  = "1"
  metric_name         = "EstimatedCharges"
  namespace           = "AWS/Billing"
  period              = "86400"
  statistic            = "Maximum"
  threshold            = var.environment == "production" ? "20" : "5"
  alarm_description   = "Lambda monthly costs exceeded threshold"
  alarm_actions       = []

  dimensions = {
    Currency    = "USD"
    ServiceName = "AWSLambda"
  }
}
```

```
tags = local.common_tags
}
```

## Cost Optimization Dashboard

- Comprehensive view of service costs and storage metrics

```
resource "aws_cloudwatch_dashboard" "cost_optimization" {
  dashboard_name = "cost-optimization-${var.environment}"

  dashboard_body = jsonencode({
    widgets = [
      {
        type    = "metric"
        width   = 12
        height  = 6
        properties = {
          metrics = [
            ["AWS/Billing", "EstimatedCharges", "ServiceName",
"AmazonS3", "Currency", "USD"],
            ["AWS/Billing", "EstimatedCharges", "ServiceName",
"AWSLambda", "Currency", "USD"],
            ["AWS/Billing", "EstimatedCharges", "ServiceName",
"AmazonRDS", "Currency", "USD"],
            ["AWS/Billing", "EstimatedCharges", "ServiceName",
"AmazonCloudFront", "Currency", "USD"],
            ["AWS/Billing", "EstimatedCharges", "ServiceName",
"AmazonApiGateway", "Currency", "USD"]
          ]
          period = 86400
          stat   = "Maximum"
          region = "us-east-1" # Billing metrics are only in us-east-1
          title  = "Service Costs (Daily)"
          yAxis = {
            left = {
              min = 0
            }
          }
        }
      },
      {
        type    = "metric"
        width   = 12
        height  = 6
        properties = {
          metrics = [
            ["AWS/S3", "BucketSizeBytes", "BucketName",
module.s3.website_bucket_name, "StorageType", "StandardStorage"],
            ["AWS/S3", "NumberOfObjects", "BucketName",
module.s3.website_bucket_name, "StorageType", "AllStorageTypes"]
          ]
        }
      }
    ]
  })
}
```

```

    ]
    period = 86400
    stat   = "Average"
    region = var.aws_region
    title  = "S3 Storage Metrics"
  }
}
]
})
}

```

## 9. Lambda Cost Optimization

File: `infra/lambda-cost-optimization.tf` - Lines 1-35

### Conditional Provisioned Concurrency

- Provisioned concurrency only enabled in production to avoid unnecessary costs

```

resource "aws_lambda_provisioned_concurrency_config"
"contact_conditional" {
  count = var.environment == "production" ? 1 : 0

  function_name           = module.lambda.lambda_function_name
  provisioned_concurrent_executions = 2
  qualifier                = "$LATEST"

  depends_on = [module.lambda]
}

```

### Cost-Optimized Lambda Logging

- Environment-based log retention for Lambda function logs

```

resource "aws_cloudwatch_log_group" "lambda_logs" {
  name =
"/aws/lambda/${module.lambda.lambda_function_name}"
  retention_in_days = var.environment == "production" ? 30 : 7
  tags              = local.common_tags
}

```

### Lambda Function URL for Development

- Direct invocation capability to reduce API Gateway costs in non-production

```

resource "aws_lambda_function_url" "contact_direct" {
  count = var.environment != "production" ? 1 : 0 # Only for non-prod
to reduce costs

  function_name      = module.lambda.lambda_function_name
  authorization_type = "NONE"

  cors {
    allow_credentials = false
    allow_origins     = ["*"]
    allow_methods     = ["POST", "OPTIONS"]
    allow_headers     = ["date", "keep-alive", "content-type"]
    expose_headers    = ["date", "keep-alive"]
    max_age           = 86400
  }
}

```

---

## 10. Module Configuration Updates

File: `infra/main.tf` - Lines 46-49

### CloudFront Module with Cost Optimizations

- Environment-based price class and log retention configuration

```

module "cloudfront" {
  source = "../modules/cloudfront"

  s3_bucket_regional_domain_name =
module.s3.website_bucket_regional_domain_name
  s3_bucket_name                  = module.s3.website_bucket_name
  price_class                     = var.environment == "production" ?
"PriceClass_All" : "PriceClass_100"
  log_retention_days              = var.environment == "production" ? 90
: 30
  tags                            = local.common_tags
}

```

---

## Cost Savings Summary

Estimated Annual Savings by Category:

1. **S3 Storage Optimization:** 20-50% reduction
  - Intelligent-Tiering: Automatic cost optimization
  - Enhanced lifecycle rules: Aggressive cleanup

- Cross-region replication optimization: 50% savings on standby storage

## 2. **CloudFront Optimization:** ~30% reduction

- PriceClass\_100 for development: Reduced global distribution costs
- Access logging with cleanup: Cost analysis capability

## 3. **RDS Multi-AZ Optimization:** ~50% reduction for non-production

- Development environments: \$240/year savings per environment
- Conditional backup retention: Additional storage cost savings

## 4. **DMS Elimination:** 100% cost elimination for non-production

- Development environments: \$360/year savings per environment
- Right-sized production instances: Additional 50% savings

## 5. **CloudWatch Log Optimization:** ~70% reduction

- Shorter retention periods: Significant log storage savings
- Environment-based configuration: Compliance with cost control

## 6. **Lambda Cost Control:** Variable savings

- Conditional provisioned concurrency: \$50-100/year savings
- Function URLs for development: API Gateway cost bypass

Total Estimated Savings:

- **Development Environment:** \$650-800/year (60-70% cost reduction)
- **Production Environment:** Maintains full functionality with 10-20% cost optimization
- **Multi-Environment Setup:** Scales linearly with environment count

All optimizations are environment-aware and automatically applied based on the **environment** variable, ensuring production maintains full reliability while development environments are aggressively cost-optimized.

## 6) Sustainability

Current state

- Static hosting with CDN and serverless compute reduces idle waste; web images include WebP formats.
- Use S3 Intelligent-Tiering for storage to optimize energy and cost based on access patterns.
- Enable S3 lifecycle rules for all buckets to automatically delete old versions and unused objects.
- Monitor and report carbon footprint using AWS CloudWatch and third-party tools (e.g., AWS Customer Carbon Footprint Tool).

Gaps

- RDS instance is always-on; caching/TTLs not optimized; no autoscaling or scheduled scale-down.
- Higher energy consumption and carbon footprint due to always-on resources and inefficient scaling.
- Missed opportunities for cost savings with S3 storage and lifecycle management.

## TF improvements

- Optimize CloudFront and S3 caching policies (set appropriate Cache-Control headers, enable Brotli/gzip compression).
- Add Terraform resources for monitoring and reporting carbon footprint (e.g., CloudWatch dashboards, custom metrics).
- Tag all resources with sustainability-related metadata (e.g., "environment", "purpose", "owner") for tracking and reporting.
- Audit and optimize static assets (images, JS, CSS) during CI/CD; automate with Terraform and build steps.

## Evidence

- `web/static/images/*.webp` (optimized images)
- `infra/modules/cloudfront/main.tf` (caching)

## Sustainability improvements (code refs)

### 1. CloudFront Compression & Caching

-Enabled Brotli/gzip compression and tiered TTLs for static assets.

```
default_cache_behavior {
  allowed_methods  = ["GET", "HEAD", "OPTIONS"]
  cached_methods  = ["GET", "HEAD", "OPTIONS"]
  target_origin_id = "s3-origin"
  compress        = true

  forwarded_values {
    query_string = false
    cookies {
      forward = "none"
    }
    headers = ["Origin", "Access-Control-Request-Method", "Access-
Control-Request-Headers"]
  }

  viewer_protocol_policy = "redirect-to-https"
  min_ttl                = 0
  default_ttl            = 86400    # 24 hours
  max_ttl                = 31536000 # 1 year
  response_headers_policy_id =
aws_cloudfront_response_headers_policy.optimized.id
}
```



## 2. S3 Lifecycle Rules

-Automatic cleanup of old versions and incomplete multipart uploads.

```
resource "aws_s3_bucket_lifecycle_configuration" "website_lifecycle" {
  bucket = aws_s3_bucket.website.id
  rule {
    id      = "cleanup_old_versions"
    status  = "Enabled"
    filter {
      prefix = ""
    }
    noncurrent_version_expiration {
      noncurrent_days = 30
    }
    abort_incomplete_multipart_upload {
      days_after_initiation = 7
    }
  }
}
```

## 3. Carbon Footprint Monitoring

-CloudWatch dashboard and custom metrics for sustainability reporting.

```
properties = {
  metrics = [
    ["AWS/Usage", "CarbonFootprint", "Service", "EC2", { stat =
"Sum" }],
    ["AWS/Usage", "CarbonFootprint", "Service", "S3", { stat =
"Sum" }],
    ["AWS/Usage", "CarbonFootprint", "Service", "Lambda", { stat
= "Sum" }]
  ]
  period = 86400
  region = data.aws_region.current.name
  title  = "AWS Carbon Footprint (Daily)"
}
```

## 4. Resource Tagging

-Sustainability-related metadata tags (environment, purpose, owner) on all resources.

```
# Common tags
locals {
```

```

common_tags = {
  Environment    = "development"
  Project        = "assignment"
  ManagedBy      = "terraform"
  Purpose        = "sustainability"
  Owner          = "DevOpsTeam"
  Sustainability = "true"
}
}

```

## 5. CI/CD Asset Optimization

-Automated image, JS, and CSS optimization in build pipeline.

-buildspec-web.yml

```

# Optimize static assets (images, JS, CSS)
- echo "Optimizing static assets..."
- npm --prefix web run optimize:assets || echo "Asset optimization
skipped (no script)"
- echo "Building static site..."

```

-package.json

```

"optimize:assets": "npx imagemin static/images/* --out-
dir=static/images && npx terser static/js/*.js -o static/js/ --compress
--mangle && npx postcss static/css/*.css -o static/css/"

```

---

## Action Items (Optional)

- Priority P0:
  - Remove RDS public ingress and move DB to private subnets with strict SGs.
  - Add API authentication (e.g., Cognito/JWT) and attach AWS WAF to API.
  - Add CloudWatch alarms + SNS for CodeBuild/CodePipeline failures and Lambda errors.
- Priority P1:
  - Place Lambda in VPC and restrict egress; add DLQ and reserved concurrency.
  - Re-enable S3 lifecycle for website/artifacts; set Cache-Control headers via deploy step.
  - Tighten IAM policies for CodeBuild/CodePipeline to least privilege.
- Priority P2:
  - Enable API Gateway access logs/caching; add dashboards for API and Lambda.
  - Evaluate Multi-AZ for RDS and Performance Insights; consider Aurora Serverless v2.
  - Add Terratest-based infra tests in CI and manual approval gates for prod.

## Notes (Optional)

-